

ARM[®] Compiler toolchain

Version 4.1

Using the Assembler



ARM Compiler toolchain

Using the Assembler

Copyright © 2010 ARM. All rights reserved.

Release Information

The following changes have been made to this book.

Change History

Date	Issue	Confidentiality	Change
May 2010	A	Non Confidential	ARM Compiler v4.1 Release

Proprietary Notice

Words and logos marked with ® or ™ are registered trademarks or trademarks of ARM® in the EU and other countries, except as otherwise stated below in this proprietary notice. Other brands and names mentioned herein may be the trademarks of their respective owners.

Neither the whole nor any part of the information contained in, or the product described in, this document may be adapted or reproduced in any material form except with the prior written permission of the copyright holder.

The product described in this document is subject to continuous developments and improvements. All particulars of the product and its use contained in this document are given by ARM in good faith. However, all warranties implied or expressed, including but not limited to implied warranties of merchantability, or fitness for purpose, are excluded.

This document is intended only to assist the reader in the use of the product. ARM shall not be liable for any loss or damage arising from the use of any information in this document, or any error or omission in such information, or any incorrect use of the product.

Where the term ARM is used it means “ARM or any of its subsidiaries as appropriate”.

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by ARM and the party that ARM delivered this document to.

Product Status

The information in this document is final, that is for a developed product.

Web Address

<http://www.arm.com>

Contents

ARM Compiler toolchain Using the Assembler

Chapter 1	Conventions and feedback	
Chapter 2	Overview of the Assembler	
2.1	About the ARM Compiler toolchain assemblers	2-2
2.2	Key features of the assembler	2-3
Chapter 3	Overview of the ARM Architecture	
3.1	About the ARM architecture	3-3
3.2	ARM, Thumb, and ThumbEE instruction sets	3-4
3.3	Changing between ARM, Thumb, and ThumbEE state	3-5
3.4	Processor modes, and privileged and unprivileged software execution	3-6
3.5	Processor modes in ARMv6-M and ARMv7-M	3-8
3.6	NEON technology	3-9
3.7	VFP coprocessor	3-10
3.8	ARM registers	3-11
3.9	General-purpose registers	3-14
3.10	Predeclared core register names	3-15
3.11	Predeclared extension register names	3-16
3.12	Predeclared XScale register names	3-17
3.13	Predeclared coprocessor names	3-19
3.14	Program Counter	3-20
3.15	Application Program Status Register	3-21

3.16	The Q flag	3-22
3.17	Current Program Status Register	3-23
3.18	Saved Program Status Registers (SPSRs)	3-24
3.19	Instruction set overview	3-25
3.20	Media processing instructions	3-27
3.21	Access to the inline barrel shifter	3-28

Chapter 4

Structure of Assembly Language Modules

4.1	Syntax of source lines in assembly language	4-2
4.2	Literals	4-4
4.3	ELF sections and the AREA directive	4-5
4.4	An example ARM assembly language module	4-6

Chapter 5

Writing ARM Assembly Language

5.1	Unified Assembler Language	5-3
5.2	Subroutines calls	5-4
5.3	Load immediates into registers	5-6
5.4	Load immediate values using MOV and MVN	5-7
5.5	Load 32-bit values to a register using MOV32	5-11
5.6	Load immediate 32-bit values to a register using LDR Rd, =const	5-12
5.7	Literal pools	5-13
5.8	Load addresses into registers	5-15
5.9	Load addresses to a register using ADR	5-16
5.10	Load addresses to a register using ADRL	5-19
5.11	Load addresses to a register using LDR Rd, =label	5-20
5.12	Other ways to Load and store registers	5-24
5.13	Load and store multiple register instructions	5-25
5.14	Load and store multiple instructions available in ARM and Thumb	5-26
5.15	Stack implementation using LDM and STM	5-28
5.16	Stack operations for nested subroutines	5-30
5.17	Block copy with LDM and STM	5-31
5.18	Read-Modify-Write procedure	5-33
5.19	Optional hash	5-34
5.20	Use of macros	5-35
5.21	Test-and-branch macro example	5-36
5.22	Unsigned integer division macro example	5-37
5.23	Symbol versions	5-39
5.24	Frame directives	5-40
5.25	Exception tables and Unwind tables	5-41
5.26	Assembly language changes after RVCTv2.1	5-42

Chapter 6

Condition Codes

6.1	Conditional instructions	6-2
6.2	Conditional execution in ARM state	6-3
6.3	Conditional execution in Thumb state	6-4
6.4	Updates to the ALU status flags	6-5

6.5	Condition code suffixes	6-7
6.6	Condition code meanings	6-9
6.7	Benefits of using conditional execution	6-11
6.8	Illustration of the benefits of using conditional instructions	6-12
6.9	Optimization for execution speed	6-16

Chapter 7 Using the Assembler

7.1	Assembler command line syntax	7-2
7.2	Assembler commands listed in groups	7-3
7.3	Specify command-line options with an environment variable	7-6
7.4	Built-in variables and constants	7-7
7.5	Versions of armasm	7-12
7.6	Diagnostic messages	7-13
7.7	Interlocks diagnostics	7-14
7.8	IT block generation	7-15
7.9	Thumb branch target alignment	7-16
7.10	Instruction width	7-17
7.11	Using the C preprocessor	7-18
7.12	Address alignment	7-20
7.13	Instruction width selection in Thumb	7-21

Chapter 8 Symbols, Literals, Expressions, and Operators

8.1	Symbol naming rules	8-3
8.2	Variables	8-4
8.3	Numeric constants	8-6
8.4	Assembly time substitution of variables	8-7
8.5	Register-relative and PC-relative expressions	8-9
8.6	Labels	8-11
8.7	Labels for PC-relative addresses	8-12
8.8	Labels for register-relative addresses	8-13
8.9	Labels for absolute addresses	8-14
8.10	Local labels	8-15
8.11	Syntax of local labels	8-17
8.12	String expressions	8-18
8.13	String literals	8-19
8.14	Numeric expressions	8-20
8.15	Numeric literals	8-21
8.16	Floating-point literals	8-23
8.17	Logical expressions	8-25
8.18	Logical literals	8-26
8.19	Unary operators	8-27
8.20	Binary operators	8-29
8.21	Multiplicative operators	8-30
8.22	String manipulation operators	8-31
8.23	Shift operators	8-32
8.24	Addition, subtraction, and logical operators	8-33
8.25	Relational operators	8-34

8.26	Boolean operators	8-35
8.27	Operator precedence	8-36
8.28	Difference between operator precedence in armasm and C	8-37

Chapter 9

NEON and VFP Programming

9.1	Architecture support for NEON and VFP	9-3
9.2	Half-precision extension	9-4
9.3	Fused Multiply-Add extension	9-5
9.4	Extension register bank mapping	9-6
9.5	NEON views of the register bank	9-8
9.6	VFP views of the extension register bank	9-9
9.7	Load values to VFP and NEON registers	9-10
9.8	Conditional execution of NEON and VFP instructions	9-11
9.9	Floating-point exceptions	9-12
9.10	NEON and VFP data types	9-14
9.11	NEON vectors	9-16
9.12	Normal NEON instructions	9-17
9.13	Long NEON instructions	9-18
9.14	Wide NEON instructions	9-19
9.15	Narrow NEON instructions	9-20
9.16	Saturating NEON instructions	9-21
9.17	NEON scalars	9-22
9.18	Extended notation	9-23
9.19	Polynomial arithmetic over {0,1}	9-24
9.20	NEON and VFP system registers	9-25
9.21	FPSCR, the floating-point status and control register	9-26
9.22	FPEXC, the floating-point exception register	9-29
9.23	FPSID, the floating-point system ID register	9-30
9.24	Flush-to-zero mode	9-31
9.25	When to use flush-to-zero mode	9-32
9.26	The effects of using flush-to-zero mode	9-33
9.27	Operations not affected by flush-to-zero mode	9-34
9.28	VFP vector mode	9-35
9.29	Vectors in the VFP extension register bank	9-36
9.30	VFP vector wrap-around	9-38
9.31	VFP vector stride	9-39
9.32	Restriction on vector length	9-40
9.33	Control of scalar, vector, and mixed operations	9-41
9.34	VFP directives and vector notation	9-43
9.35	Pre-UAL VFP mnemonics	9-44
9.36	Vector notation	9-46
9.37	VFPASSERT SCALAR	9-47
9.38	VFPASSERT VECTOR	9-48

Chapter 1

Conventions and feedback

The following describes the typographical conventions and how to give feedback:

Typographical conventions

The following typographical conventions are used:

- `monospace` Denotes text that can be entered at the keyboard, such as commands, file and program names, and source code.
- `monospace` Denotes a permitted abbreviation for a command or option. The underlined text can be entered instead of the full command or option name.
- `monospace`* *italic* Denotes arguments to commands and functions where the argument is to be replaced by a specific value.
- `monospace` **bold** Denotes language keywords when used outside example code.
- italic* Highlights important notes, introduces special terminology, denotes internal cross-references, and citations.
- bold** Highlights interface elements, such as menu names. Also used for emphasis in descriptive lists, where appropriate, and for ARM® processor signal names.

Feedback on this product

If you have any comments and suggestions about this product, contact your supplier and give:

- your name and company
- the serial number of the product
- details of the release you are using
- details of the platform you are using, such as the hardware platform, operating system type and version
- a small standalone sample of code that reproduces the problem
- a clear explanation of what you expected to happen, and what actually happened
- the commands you used, including any command-line options
- sample output illustrating the problem
- the version string of the tools, including the version number and build numbers.

Feedback on content

If you have comments on content then send an e-mail to errata@arm.com. Give:

- the title
- the number, ARM DUI 0473A
- if viewing online, the topic names to which your comments apply
- if viewing a PDF version of a document, the page numbers to which your comments apply
- a concise explanation of your comments.

ARM also welcomes general suggestions for additions and improvements.

ARM periodically provides updates and corrections to its documentation on the ARM Information Center, together with knowledge articles and *Frequently Asked Questions* (FAQs).

Other information

- ARM Information Center, <http://infocenter.arm.com/help/index.jsp>
- ARM Technical Support Knowledge Articles, <http://infocenter.arm.com/help/topic/com.arm.doc.faq/index.html>
- Keil Distributors, <http://www.keil.com/distis>.

Chapter 2

Overview of the Assembler

This following topics introduce the assemblers provided with ARM Compiler toolchain:

- *About the ARM Compiler toolchain assemblers* on page 2-2
- *Key features of the assembler* on page 2-3.

2.1 About the ARM Compiler toolchain assemblers

ARM Compiler toolchain provides:

- A freestanding assembler, `armasm`.
- An optimizing inline assembler and a non-optimizing embedded assembler built into the C and C++ compilers. These use the same syntax for assembly instructions.

2.1.1 See also

Concepts

Developing for ARM Processors:

- Chapter 4 *Mixing C, C++, and Assembly Language*.

Using the Compiler:

- Chapter 8 *Using the Inline and Embedded Assemblers of the ARM Compiler*.

Migration and Compatibility:

- Chapter 3 *Migrating from RVCT v4.0 to ARM Compiler v4.1*
- Chapter 4 *Migrating from RVCT v3.1 to RVCT v4.0*.

2.2 Key features of the assembler

The ARM® assembler supports:

- *Unified Assembly Language* (UAL) for both ARM and Thumb® code
- *NEON Single Instruction Multiple Data* (SIMD) instructions in ARM and Thumb code
- *Vector Floating Point* (VFP) instructions in ARM and Thumb code
- *Wireless MMX Technology* instructions to assemble code to run on the PXA270 processor
- directives in assembly source code
- processing of user defined macros.

2.2.1 See also

Concepts

- *Unified Assembler Language* on page 5-3
- *NEON technology* on page 3-9
- *Use of macros* on page 5-35
- *Architecture support for NEON and VFP* on page 9-3.

Reference

Assembler Reference:

- Chapter 4 *NEON and VFP Programming*
- Chapter 5 *Wireless MMX Technology Instructions*
- Chapter 6 *Directives Reference*.

Chapter 3

Overview of the ARM Architecture

The following topics give an overview of the ARM architecture:

- *About the ARM architecture* on page 3-3
- *ARM, Thumb, and ThumbEE instruction sets* on page 3-4
- *Changing between ARM, Thumb, and ThumbEE state* on page 3-5
- *Processor modes, and privileged and unprivileged software execution* on page 3-6
- *Processor modes in ARMv6-M and ARMv7-M* on page 3-8
- *NEON technology* on page 3-9
- *VFP coprocessor* on page 3-10
- *ARM registers* on page 3-11
- *General-purpose registers* on page 3-14
- *Predeclared core register names* on page 3-15
- *Predeclared extension register names* on page 3-16
- *Predeclared XScale register names* on page 3-17
- *Predeclared coprocessor names* on page 3-19
- *Program Counter* on page 3-20
- *Application Program Status Register* on page 3-21
- *The Q flag* on page 3-22

- *Current Program Status Register* on page 3-23
- *Saved Program Status Registers (SPSRs)* on page 3-24
- *Instruction set overview* on page 3-25
- *Media processing instructions* on page 3-27
- *Access to the inline barrel shifter* on page 3-28.

3.1 About the ARM architecture

ARM processors are typical of RISC processors in that they implement a load and store architecture. Only load and store instructions can access memory. Data processing instructions operate on register contents only.

It is assumed that you are using a processor that implements the ARMv4 or later architecture. All these processors have a 32-bit addressing range.

3.1.1 See also

Reference

- *ARM Architecture Reference Manual*,
<http://infocenter.arm.com/help/topic/com.arm.doc.subset.arch.reference/index.html>.

3.2 ARM, Thumb, and ThumbEE instruction sets

The ARM instruction set is a set of 32-bit instructions providing a comprehensive range of operations.

ARMv4T and later define a 16-bit instruction set called the Thumb instruction set. Most of the functionality of the 32-bit ARM instruction set is available, but some operations require more instructions. The Thumb instruction set provides better code density, at the expense of performance.

ARMv6T2 introduces a major enhancement of the Thumb instruction set by providing 32-bit Thumb instructions. The 32-bit and 16-bit Thumb instructions together provide almost exactly the same functionality as the ARM instruction set. The enhanced Thumb instruction set (Thumb-2) achieves the high performance of ARM code and better code density like 16-bit Thumb code.

ARMv7 defines the Thumb Execution Environment (ThumbEE). The ThumbEE instruction set is based on Thumb, with some changes and additions to make it a better target for dynamically generated code, that is, code compiled on the device either shortly before or during execution.

3.2.1 See also

Concepts

- *Instruction set overview* on page 3-25.

3.3 Changing between ARM, Thumb, and ThumbEE state

A processor that is executing ARM instructions is operating in *ARM state*. A processor that is executing Thumb instructions is operating in *Thumb state*. A processor that is executing ThumbEE instructions is operating in *ThumbEE state*. A processor can also operate in another state called the *Jazelle state*. The assembler cannot directly assemble code for the Jazelle state.

Each instruction set includes instructions to change processor state.

To change between ARM and Thumb states, you must switch the assembler mode to produce the correct opcodes using ARM or THUMB directives. To generate ThumbEE code, use THUMBX. Assembler code using CODE32 and CODE16 can still be assembled by the assembler, but you are recommended to use ARM and THUMB for new code.

A processor in one state cannot execute instructions from another instruction set. For example, a processor in ARM state cannot execute Thumb instructions, and a processor in Thumb state cannot execute ARM instructions. You must ensure that the processor never receives instructions of the wrong instruction set for the current state.

The initial state after reset depends on the processor being used and its configuration.

3.3.1 See also

Reference

Assembler Reference:

- *B, BL, BX, BLX, and BXJ* on page 3-130
- *ENTERX and LEAVEX* on page 3-174
- *Instruction set and syntax selection directives* on page 6-74.

3.4 Processor modes, and privileged and unprivileged software execution

ARM processors support different processor modes depending on the architecture version. See Table 3-1.

———— **Note** ————

ARMv6-M and ARMv7-M do not support the same modes as other ARM processors. The processor modes described here do not apply to ARMv6-M and ARMv7-M.

Table 3-1 ARM processor modes

Processor mode	Architectures	Mode number
User	All	0b10000
FIQ - Fast Interrupt Request	All	0b10001
IRQ - Interrupt Request	All	0b10010
Supervisor	All	0b10011
Abort	All	0b10111
Undefined	All	0b11011
System	ARMv4 and later	0b11111
Monitor	Security Extensions only	0b10110

User mode is an unprivileged mode, and has restricted access to system resources. All other modes have full access to system resources in the current security state, can change mode freely, and execute software as privileged.

Applications that require task protection usually execute in User mode. Some embedded applications might run entirely in any mode other than User mode. An application that requires full access to system resources usually executes in System mode.

Modes other than User mode are entered to service exceptions, or to access privileged resources.

On an implementation that includes the Security Extensions, code can run in either a secure state or in a non-secure state.

3.4.1 See also

Concepts

- *Processor modes in ARMv6-M and ARMv7-M* on page 3-8.

Developing code for ARM Processors:

- Chapter 6 *Handling Processor Exceptions*.

Reference

- *ARM Architecture Reference Manual*,
<http://infocenter.arm.com/help/topic/com.arm.doc.subset.arch.reference/index.html>.

3.5 Processor modes in ARMv6-M and ARMv7-M

In ARMv6-M and ARMv7-M there are two processor modes available:

- Thread mode
- Handler mode.

Thread mode is the normal mode that programs run in. Thread mode can be privileged or unprivileged software execution. Handler mode is the mode that exceptions are handled in. The handler mode is always privileged software execution.

3.5.1 See also

Concepts

- *Processor modes, and privileged and unprivileged software execution* on page 3-6.

Developing Software for ARM Processors:

- Chapter 6 *Handling Processor Exceptions*.

3.6 NEON technology

ARM NEON™ Technology is the implementation of the Advanced SIMD architecture extension. It is a 64 and 128-bit hybrid SIMD technology targeted at advanced media and signal processing applications and embedded processors. It is implemented as part of the ARM core, but has its own execution pipelines and a register bank that is distinct from the ARM core register bank.

NEON instructions are available in both ARM and Thumb code.

3.6.1 See also

Concepts

- *Architecture support for NEON and VFP* on page 9-3
- *NEON views of the register bank* on page 9-8
- *Chapter 9 NEON and VFP Programming*
- *Chapter 4 Using the NEON Vectorizing Compiler.*

3.7 VFP coprocessor

The VFP coprocessor, together with associated support code, provides single-precision and double-precision floating-point arithmetic, as defined by *ANSI/IEEE Std. 754-1985 IEEE Standard for Binary Floating-Point Arithmetic*. This document is referred to as the IEEE 754 standard.

The VFP coprocessor uses a register bank that is distinct from the ARM core register bank.

Note

The VFP register bank is shared with the NEON register bank.

Your processor might have either the VFPv2, VFPv3, or VFPv4 coprocessor. There are variants of VFPv3 that differ in the number of accessible registers or differ in its support of the half-precision extension:

- VFPv3
- VFPv3-D16
- VFPv3-FP16
- VFPv3-D16-FP16.

3.7.1 See also

Concepts

- *Architecture support for NEON and VFP* on page 9-3
- *Half-precision extension* on page 9-4
- *VFP views of the extension register bank* on page 9-9
- Chapter 9 *NEON and VFP Programming*.

3.8 ARM registers

In all ARM processors, the following registers are available and accessible in any processor mode:

- 13 general-purpose registers R0-R12
- 1 *Stack Pointer* (SP)
- 1 *Link Register* (LR)
- 1 *Program Counter* (PC)
- 1 *Application Program Status Register* (APSR)

Note

- The Link Register can also be used as a general-purpose register. The Stack Pointer can be used as a general-purpose register in ARM state only.
-

Additional registers are available in privileged software execution. ARM processors, with the exception of ARMv6-M and ARMv7-M based processors, have a total of 37 or 40 registers depending on whether the Security Extensions are implemented. The registers are arranged in partially overlapping banks. There is a different register bank for each processor mode. The banked registers give rapid context switching for dealing with processor exceptions and privileged operations.

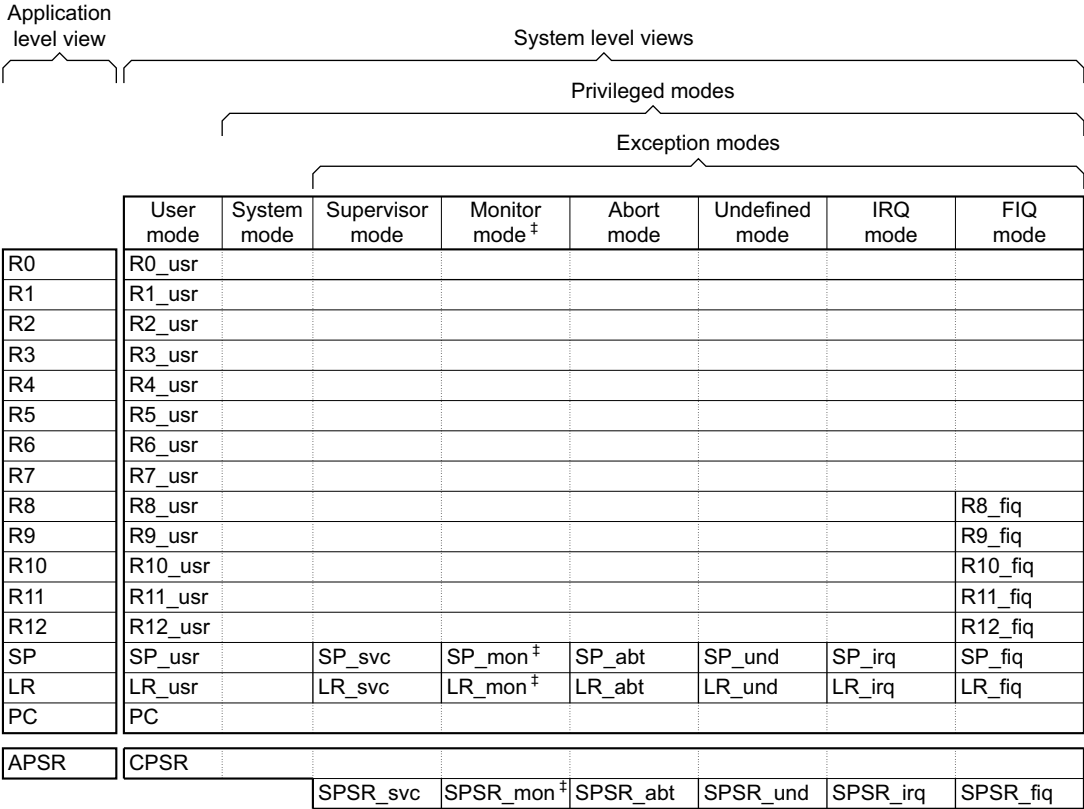
The additional registers in ARM processors, with the exception of ARMv6-M and ARMv7-M, are:

- 2 supervisor mode registers for banked SP and LR
- 2 abort mode registers for banked SP and LR
- 2 undefined mode registers for banked SP and LR
- 2 interrupt mode registers for banked SP and LR
- 7 FIQ mode registers for banked R8-R12, SP and LR
- 2 monitor mode registers for banked SP and LR
- 6 *Saved Program Status Register* (SPSRs), one for each exception mode.

Note

- The monitor mode registers and one of the SPSRs apply only to the monitor mode and are only present if Security Extensions are implemented.
 - In privileged software execution, CPSR is an alias for APSR and gives access to additional bits.
-

Figure 3-1 on page 3-12 shows how the registers are banked in the ARM Architecture except ARMv6-M and ARMv7-M.



‡ Monitor mode and the associated banked registers are implemented only as part of the Security Extensions

Figure 3-1 Organization of general-purpose registers and Program Status Registers

In ARMv6-M and ARMv-7M based processors, SP is an alias for the two banked stack pointer registers:

- Main stack pointer register, which is only available in privileged software execution.
- Process stack pointer register.

3.8.1 See also

Concepts

- *General-purpose registers* on page 3-14
- *Program Counter* on page 3-20
- *Application Program Status Register* on page 3-21
- *Saved Program Status Registers (SPSRs)* on page 3-24

- *Current Program Status Register* on page 3-23
- *Processor modes, and privileged and unprivileged software execution* on page 3-6.

Reference

- *ARM Architecture Reference Manual*,
<http://infocenter.arm.com/help/topic/com.arm.doc.subset.arch.reference/index.html>.

3.9 General-purpose registers

With the exception of ARMv6-M and ARMv7-M based processors, there are 30 (or 32 if Security Extensions are implemented) general-purpose 32-bit registers, that include the banked SP and LR registers. Fifteen general-purpose registers are visible at any one time, depending on the current processor mode. These are R0-R12, SP, LR.

SP (or R13) is the *stack pointer*. The C and C++ compilers always use SP as the stack pointer. In Thumb SP is strictly defined as the stack pointer. So, apart from a few designated stack manipulation instructions, most instructions are unpredictable if they use SP. Use of SP as a general purpose register is discouraged.

In User mode, LR (or R14) is used as a *link register* to store the return address when a subroutine call is made. It can also be used as a general-purpose register if the return address is stored on the stack.

In the exception handling modes, LR holds the return address for the exception, or a subroutine return address if subroutine calls are executed within an exception. LR can be used as a general-purpose register if the return address is stored on the stack.

In ARM state, all instructions can access R0 to R12, SP, and LR, and most instructions can also access PC (R15). The MRS instructions can move the contents of a status register to a general-purpose register, where they can be manipulated by normal data processing operations. The MSR instruction can be used to move the contents of a general-purpose register to a status register.

Most 16-bit Thumb instructions can only access R0 to R7. Only a small number of these instructions can access R8-R12, SP, LR, and PC. Registers R0 to R7 are called Lo registers. Registers R8-R12, SP, LR, and PC are called Hi registers.

Regarding register accesses, 32-bit Thumb instructions provide the same facilities as ARM instructions, except that some of the less useful accesses to SP and PC are not permitted in 32-bit Thumb.

3.9.1 See also

Concepts

- *Program Counter* on page 3-20
- *Predeclared core register names* on page 3-15.

Reference

Assembler Reference:

- *MRS* on page 3-155
- *MSR* on page 3-157.

3.10 Predeclared core register names

Table 3-2 shows the predeclared core registers:

Table 3-2 Predeclared core registers

Register names	Meaning
r0-r15 and R0-R15	General purpose registers.
a1-a4	Argument, result or scratch registers. These are synonyms for R0 to R3
v1-v8	Variable registers. These are synonyms for R4 to R11.
sb and SB	Static base register. This is a synonym for R9.
ip and IP	Intra procedure call scratch register. This is a synonym for R12.
sp and SP	Stack pointer. This is a synonym for R13.
lr and LR	Link register. This is a synonym for R14.
pc and PC	Program counter. This is a synonym for R15.

With the exception of a1-a4 and v1-v8, you can write the registers either in all upper case or all lower case.

3.10.1 See also

Concepts

- *General-purpose registers* on page 3-14.

3.11 Predeclared extension register names

The following extension register names are predeclared:

Table 3-3 Predeclared extension registers

Register names	Meaning
q0-q15 and Q0-Q15	NEON quadword registers.
d0-d31 and D0-D31	NEON doubleword registers, VFP double-precision registers.
s0-s31 and S0-S31	VFP single-precision registers.

You can write the registers either in upper case or lower case.

3.11.1 See also

Concepts

- *Extension register bank mapping* on page 9-6.

3.12 Predeclared XScale register names

The following register names are predeclared when assembling for a Marvell XScale CPU:

Table 3-4 Predeclared XScale registers

Register name	Meaning
acc0-acc7 and ACC0-ACC7	XScale accumulator registers.

The following register names are predeclared when assembling for a Marvell XScale CPU with Wireless MMX:

Table 3-5 Predeclared Wireless MMX registers

Register name	Meaning
wR0-wR15, wr0-wr15, and WR0-WR15	Wireless SIMD data registers (coprocessor 0).
wC0-wC15, wc0-wc15, and WC0-WC15	Usable aliases for coprocessor 1 registers. Use of these aliases is not recommended.
wCID, wcid, and WCID	Coprocessor ID register (coprocessor 1 register c0).
wCon, wcon, and WCON	Control register (coprocessor 1 register c1).
wCSSF, wcssf, and WCSSF	Saturation SIMD flags (coprocessor 1 register c2).
wCASF, wcasf, and WCASF	Arithmetic SIMD flags (coprocessor 1 register c3).
wCGR0-wCGR3, wgr0-wgr3, and WCGR0-WCGR3	General purpose registers (coprocessor 1 registers c8 - c11).

The register names are case-sensitive and can be mixed case where this matches exactly the Wireless MMX Technology specification.

Control registers, ID register, general-purpose registers wCGR0 - wCGR3 and the SIMD flags map onto coprocessor 1. Use the Wireless MMX Technology instructions TMCR and TMRC to read and write to these registers. The coprocessor 1 registers c4-c7 and c12-c15 are reserved.

SIMD data registers (wR0 - wR15) map onto coprocessor 0 and hold 16x64-bit packed data. Use the Wireless MMX Technology pseudo-instructions TMRRC and TMCRR to move data between these registers and the ARM registers.

The assembler supports the WRN and WCN directives to specify your own register names.

3.12.1 See also

Reference

Assembler Reference:

- *ARM support for Wireless MMX technology* on page 5-3
- *Wireless MMX pseudo-instructions* on page 5-9.

Wireless MMX Technology Developer Guide.

3.13 Predeclared coprocessor names

The following coprocessor names and coprocessor register names are predeclared:

Table 3-6 Predeclared coprocessor registers

Register name	Meaning
p0-p15	Coprocessors 0-15.
c0-c15	Coprocessor registers 0-15.

All register and coprocessor names are case-sensitive.

3.13.1 See also

Reference

Assembler Reference:

- *CDP and CDP2* on page 3-140
- *MCR, MCR2, MCRR, and MCRR2* on page 3-141
- *MRC, MRC2, MRRC and MRRC2* on page 3-143.

3.14 Program Counter

The *Program Counter* (PC) is accessed as PC (or R15). It is incremented by the size of the instruction executed (which is always four bytes in ARM state). Branch instructions load the destination address into PC. You can also load the PC directly using data operation instructions. For example, to branch to the address in a general purpose register, use:

```
MOV PC,R0
```

During execution, PC does not contain the address of the currently executing instruction. The address of the currently executing instruction is typically PC-8 for ARM, or PC-4 for Thumb.

———— **Note** ————

Writing to the PC directly is not the recommended method for jumping to an address or returning from a function. Use the BX instruction instead.

—————

3.14.1 See also

Concepts

- *Register-relative and PC-relative expressions* on page 8-9.

Reference

- *B, BL, BX, BLX, and BXJ* on page 3-130.

3.15 Application Program Status Register

The *Application Program Status Register* (APSR) holds copies of the *Arithmetic Logic Unit* (ALU) status flags. They are also known as the condition code flags. They are used to determine whether conditional instructions are executed or not.

On ARMv5TE, ARMv6 and later architectures, the APSR also holds the Q (saturation) flag.

On ARMv6 and later, the APSR also holds the GE (Greater than or Equal) flags. The GE flags can be set by the parallel add and subtract instructions. The GE flags are used by the SEL instruction to perform byte-based selection from two registers.

These flags are accessible in all modes, using MSR and MRS instructions.

3.15.1 See also

Concepts

- *Updates to the ALU status flags* on page 6-5
- *Conditional instructions* on page 6-2.

Reference

Assembler Reference:

- *MRS* on page 3-155
- *MSR* on page 3-157
- *SEL* on page 3-77
- *Parallel add and subtract* on page 3-113.

3.16 The Q flag

ARMv5TE, ARMv6 and later have a Q flag that is set to 1 when saturation has occurred in saturating arithmetic instructions, or when overflow has occurred in certain multiply instructions.

The Q flag is a *sticky* flag. Although the saturating and certain multiply instructions can set the flag, they cannot clear it. You can execute a series of such instructions, and then test the flag to find out whether saturation or overflow occurred at any point in the series, without having to check the flag after each instruction.

To clear the Q flag, use an MSR instruction to read-modify-write the APSR:

```
MRS r5, APSR
BIC r5, r5, #(1<<27)
MSR APSR_nzcvq, r5
```

The state of the Q flag cannot be tested directly by the condition codes. To read the state of the Q flag, use an MRS instruction.

```
MRS r6, APSR
TST r6, #(1<<27); Z is clear if Q flag was set
```

3.16.1 See also

Concepts

- *Read-Modify-Write procedure* on page 5-33.

Reference

Assembler Reference:

- *MRS* on page 3-155
- *MSR* on page 3-157
- *QADD, QSUB, QDADD, and QDSUB* on page 3-108
- *SMULxy and SMLAxy* on page 3-90
- *SMULWy and SMLAWy* on page 3-92.

3.17 Current Program Status Register

The *Current Program Status Register* (CPSR) holds:

- the APSR flags
- the current processor mode
- interrupt disable flags
- current processor state (ARM, Thumb, ThumbEE, or Jazelle®)
- endianness state (on ARMv4T and later)
- execution state bits for the IT block (on ARMv6T2 and later).

The execution state bits control conditional execution in the IT block.

Only the APSR flags are accessible in all modes. The endianness bit (E) of the CPSR is accessible only in privileged software execution. It can be read by MRS and written by MSR, but SETEND is the preferred instruction to write to the E bit.

The execution state bits for the IT block (IT[1:0]), Jazelle bit (J), and Thumb bit (T) can be accessed by MRS only in Debug state.

3.17.1 See also

Concepts

- *Updates to the ALU status flags* on page 6-5
- *Saved Program Status Registers (SPSRs)* on page 3-24.

Reference

Assembler Reference:

- *IT* on page 3-134
- *SETEND* on page 3-163
- *MSR* on page 3-157
- *MRS* on page 3-155.

3.18 Saved Program Status Registers (SPSRs)

The SPSR is used to store the current value of the CPSR when an exception is taken so that it can be restored after handling the exception. Each exception handling mode can access its own SPSR. User mode and System mode do not have an SPSR because they are not exception handling modes.

The execution state bits, endianness state and current processor state can be accessed from the SPSR in any exception mode, using the MSR and MRS instruction.

3.18.1 See also

Concepts

- *Current Program Status Register* on page 3-23.
- Developing code for ARM Processors:*
- Chapter 6 *Handling Processor Exceptions*.

3.19 Instruction set overview

All ARM instructions are 32 bits long. Instructions are stored word-aligned, so the least significant two bits of instruction addresses are always zero in ARM state.

Thumb and ThumbEE instructions are either 16 or 32 bits long. Instructions are stored half-word aligned. Some instructions use the least significant bit of the address to determine whether the code being branched to is Thumb code or ARM code.

Before the introduction of Thumb-2, the Thumb instruction set was limited to a restricted subset of the functionality of the ARM instruction set. Almost all Thumb instructions were 16-bit. The Thumb-2 instruction set functionality is almost identical to that of the ARM instruction set.

Table 3-7 describes some of the functional groupings of the available instructions.

Table 3-7 Instruction groups

Instruction Group	Description
Branch and control	<p>These instructions are used to:</p> <ul style="list-style-type: none"> • branch to subroutines • branch backwards to form loops • branch forward in conditional structures • make following instructions conditional without branching • change the processor between ARM state and Thumb state.
Data processing	<p>These instructions operate on the general-purpose registers. They can perform operations such as addition, subtraction, or bitwise logic on the contents of two registers and place the result in a third register. They can also operate on the value in a single register, or on a value in a register and an immediate value supplied within the instruction.</p> <p>Long multiply instructions give a 64-bit result in two registers.</p>
Register load and store	<p>These instructions load or store the value of a single register from or to memory. They can load or store a 32-bit word, a 16-bit halfword, or an 8-bit unsigned byte. Byte and halfword loads can either be sign extended or zero extended to fill the 32-bit register.</p> <p>A few instructions are also defined that can load or store 64-bit doubleword values into two 32-bit registers.</p>

Table 3-7 Instruction groups (continued)

Instruction Group	Description
Multiple register load and store	These instructions load or store any subset of the general-purpose registers from or to memory.
Status register access	These instructions move the contents of a status register to or from a general-purpose register.
Coprocessor	These instructions support a general way to extend the ARM architecture. They also enable the control of the CP15 System Control coprocessor registers.

3.19.1 See also

Concepts

- *Load and store multiple register instructions* on page 5-25.

Reference

Assembler Reference:

- Chapter 3 *ARM and Thumb Instructions*.

3.20 Media processing instructions

Media processing instructions are available in the media extension for ARMv6 and later. The media processing instructions are:

Parallel add and subtract instructions:

- *Parallel add and subtract* on page 3-113.

Extend instructions:

- *SXT, SXTA, UXT, and UXTA* on page 3-124.

Multiply instructions:

- *SMLAD and SMLSD* on page 3-100
- *SMLALD and SMLSLD* on page 3-102
- *SMMUL, SMMLA, and SMMLS* on page 3-98
- *SMUAD{X} and SMUSD{X}* on page 3-96.

Packing, unpacking, saturation, and reversal instructions:

- *PKHBT and PKHTB* on page 3-127
- *REV, REV16, REVSH, and RBIT* on page 3-79
- *SEL* on page 3-77
- *SSAT and USAT* on page 3-110
- *SSAT16 and USAT16* on page 3-118.

Absolute sum and bit field instructions:

- *BFC and BFI* on page 3-121
- *SBFX and UBFX* on page 3-123
- *USAD8 and USADA8* on page 3-116.

Note

BFC and BFI are available on ARMv6T2 and above.

3.21 Access to the inline barrel shifter

The ARM arithmetic logic unit has a 32-bit barrel shifter that is capable of shift and rotate operations. The second operand to many ARM and Thumb data-processing and single register data-transfer instructions can be shifted, before the data-processing or data-transfer is executed, as part of the instruction. This supports, but is not limited to:

- scaled addressing
- multiplication by an immediate value
- constructing immediate values.

32-bit Thumb instructions give almost the same access to the barrel shifter as ARM instructions.

The 16-bit Thumb instructions only allow access to the barrel shifter using separate instructions.

3.21.1 See also

Concepts

- *Load immediates into registers* on page 5-6
- *Load immediate values using MOV and MVN* on page 5-7.

Chapter 4

Structure of Assembly Language Modules

Assembly language is the language that the assembler (armasm) parses and assembles to produce object code. The following topics describe the structure of the assembly source files:

- *Syntax of source lines in assembly language* on page 4-2
- *Literals* on page 4-4
- *ELF sections and the AREA directive* on page 4-5
- *An example ARM assembly language module* on page 4-6.

4.1 Syntax of source lines in assembly language

The general form of source lines in assembly language is:

```
{symbol} {instruction|directive|pseudo-instruction} {;comment}
```

All three sections of the source line are optional.

symbol is usually a label. In instructions and pseudo-instructions it is always a label. In some directives it is a symbol for a variable or a constant. The description of the directive makes this clear in each case.

symbol must begin in the first column. It cannot contain any white space character such as a space or a tab unless it is enclosed by bars (|).

Labels are symbolic representations of addresses. You can use labels to mark specific addresses that you want to refer to from other parts of the code. Local labels are a subclass of labels that begin with a number in the range 0-99. Unlike other labels, a local label can be defined many times. This makes them useful when generating labels with a macro.

Directives provide important information to the assembler that either affects the assembly process or affects the final output image.

Instructions and pseudo-instructions make up the code a processor uses to perform tasks.

————— **Note** —————

Instructions, pseudo-instructions, and directives must be preceded by white space, such as a space or a tab, irrespective of whether there is a preceding label or not.

Some directives do not allow the use of a label.

A comment is the final part of a source line. The first semicolon on a line marks the beginning of a comment except where the semicolon appears inside a string literal. The end of the line is the end of the comment. A comment alone is a valid line. The assembler ignores all comments. You can use blank lines to make your code more readable.

Instruction mnemonics, pseudo-instructions, directives, and symbolic register names (except a1-a4, v1-v8, and Wireless MMX registers) can be written in all uppercase or all lowercase, but not mixed. Labels and comments can be in uppercase or lowercase, or mixed.

Example 4-1

```

        AREA    ARMex, CODE, READONLY
        ENTRY                      ; Name this block of code ARMex
                                     ; Mark first instruction to execute
start
        MOV     r0, #10             ; Set up parameters
        MOV     r1, #3
        ADD     r0, r0, r1          ; r0 = r0 + r1
stop
        MOV     r0, #0x18           ; angel_SWIreason_ReportException
        LDR     r1, =0x20026         ; ADP_Stopped_ApplicationExit
        SVC     #0x123456           ; ARM semihosting (formerly SWI)
        END                      ; Mark end of file

```

To make source files easier to read, a long line of source can be split onto several lines by placing a backslash character (\) at the end of the line. The backslash must not be followed by any other characters (including spaces and tabs). The assembler treats the backslash followed by end-of-line sequence as white space. You can also use blank lines to make your code more readable.

Note

Do not use the backslash followed by end-of-line sequence within quoted strings.

The limit on the length of lines, including any extensions using backslashes, is 4095 characters.

4.1.1 See also**Concepts**

- *Labels* on page 8-11
- *Local labels* on page 8-15
- *Symbol naming rules* on page 8-3
- *Numeric literals* on page 8-21
- *String literals* on page 8-19
- *Literals* on page 4-4.

4.2 Literals

In assembly source files, literals can be expressed as:

- decimal numbers, for example 123
- hexadecimal numbers, for example 0x7B
- numbers in any base from 2 to 9, for example 5_204 is a number in base 5
- floating point numbers, for example 123.4
- boolean values {TRUE} or {FALSE}
- single character values enclosed by single quotes, for example 'w'
- strings enclosed in double quotes, for example "This is a string".

Note

In most cases, a string containing a single character is accepted as a single character value. For example `ADD r0,r1,#"a"` is accepted, but `ADD r0,r1,#"ab"` is faulted.

You can also use variables and names to represent the literals.

4.2.1 See also

Concepts

- *Syntax of source lines in assembly language* on page 4-2.

4.3 ELF sections and the AREA directive

ELF *sections* are independent, named, indivisible sequences of code or data. A single code section is the minimum required to produce an application.

The output of an assembly or compilation can include:

- one or more code sections. These are usually read-only sections.
- one or more data sections. These are usually read-write sections. They might be *zero initialized (ZI)*.

The linker places each section in a program image according to section placement rules. Sections that are adjacent in source files are not necessarily adjacent in the application image.

In a source file, the AREA directive marks the start of a section. This directive names the section and sets its attributes. The attributes are placed after the name, separated by commas.

You can choose any name for your sections. However, names starting with any nonalphabetic character must be enclosed in bars, or an AREA name missing error is generated. For example, |1_DataArea|.

Example 4-2 defines a single read-only section called ARMex that contains code.

Example 4-2

```
AREA    ARMex, CODE, READONLY
        ; Name this block of code ARMex
```

4.3.1 See also

Concepts

- *An example ARM assembly language module* on page 4-6.

Using the Linker:

- Chapter 8 *Using scatter-loading description files*.

Reference

Assembler Reference:

- *AREA* on page 6-81.

4.4 An example ARM assembly language module

Example 4-3 illustrates some of the core constituents of an assembly language module. The example is written in ARM assembly language.

The constituent parts of this example are:

- ELF sections (defined by the AREA directive)
- application entry (defined by the ENTRY directive)
- application execution
- application termination
- program end (defined by the END directive).

Example 4-3 defines a single section called ARMex that contains code and is marked as being READONLY.

Example 4-3 Constituents of an assembly language module

	AREA	ARMex, CODE, READONLY	
			; Name this block of code ARMex
	ENTRY		; Mark first instruction to execute
start	MOV	r0, #10	; Set up parameters
	MOV	r1, #3	
	ADD	r0, r0, r1	; r0 = r0 + r1
stop	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
	END		; Mark end of file

4.4.1 Application entry

The ENTRY directive declares an entry point to the program. It marks the first instruction to be executed. In applications using the C library, an entry point is also contained within the C library initialization code. Initialization code and exception handlers also contain entry points.

4.4.2 Application execution

The application code in Example 4-3 begins executing at the label start, where it loads the decimal values 10 and 3 into registers R0 and R1. These registers are added together and the result placed in R0.

4.4.3 Application termination

After executing the main code, the application terminates by returning control to the debugger. This is done using the ARM semihosting SVC (0x123456 by default), with the following parameters:

- R0 equal to `angel_SWIreason_ReportException` (0x18)
- R1 equal to `ADP_Stopped_ApplicationExit` (0x20026).

4.4.4 Program end

The `END` directive instructs the assembler to stop processing this source file. Every assembly language source module must finish with an `END` directive on a line by itself. Any lines following the `END` directive are ignored by the assembler.

4.4.5 See also

Concepts

- *ELF sections and the `AREA` directive* on page 4-5.

Developing Software for ARM processors:

- Chapter 8 *Semihosting*.

Reference

Assembler Reference:

- *ENTRY* on page 6-87
- *END* on page 6-87.

Chapter 5

Writing ARM Assembly Language

The following topics describe the use of a few basic assembler instructions and the use of macros:

- *Unified Assembler Language* on page 5-3
- *Subroutines calls* on page 5-4
- *Load immediates into registers* on page 5-6
- *Load immediate values using MOV and MVN* on page 5-7
- *Load 32-bit values to a register using MOV32* on page 5-11
- *Load immediate 32-bit values to a register using LDR Rd, =const* on page 5-12
- *Load addresses into registers* on page 5-15
- *Load addresses to a register using ADR* on page 5-16
- *Load addresses to a register using ADRL* on page 5-19
- *Load addresses to a register using LDR Rd, =label* on page 5-20
- *Other ways to Load and store registers* on page 5-24
- *Load and store multiple register instructions* on page 5-25
- *Load and store multiple instructions available in ARM and Thumb* on page 5-26
- *Stack implementation using LDM and STM* on page 5-28
- *Stack operations for nested subroutines* on page 5-30
- *Block copy with LDM and STM* on page 5-31

- *Read-Modify-Write procedure* on page 5-33
- *Optional hash* on page 5-34
- *Use of macros* on page 5-35
- *Test-and-branch macro example* on page 5-36
- *Unsigned integer division macro example* on page 5-37
- *Symbol versions* on page 5-39
- *Frame directives* on page 5-40
- *Exception tables and Unwind tables* on page 5-41
- *Assembly language changes after RVCTv2.1* on page 5-42.

5.1 Unified Assembler Language

Unified Assembler Language (UAL) is a common syntax for ARM and Thumb instructions. It supersedes earlier versions of both the ARM and Thumb assembler languages.

Code written using UAL can be assembled for ARM or Thumb for any ARM processor. The assembler faults the use of unavailable instructions.

RealView Compilation Tools (RVCT) v2.1 and earlier can only assemble the pre-UAL syntax. Later versions of RVCT and ARM Compiler toolchain can assemble code written in pre-UAL and UAL syntax.

By default, the assembler expects source code to be written in UAL. The assembler accepts UAL syntax if any of the directives `CODE32`, `ARM`, `THUMB`, or `THUMBX` is used or if you assemble with any of the `--32`, `--arm`, `--thumb`, or `--thumbx` command-line options. The assembler also accepts source code written in pre-UAL ARM assembly language when you assemble with `CODE32` or `ARM`.

The assembler accepts source code written in pre-UAL Thumb assembly language when you assemble using the `--16` command-line option, or the `CODE16` directive in the source code.

Note

The pre-UAL Thumb assembly language does not support Thumb-2 instructions.

5.1.1 See also

Concepts

- *Assembly language changes after RVCTv2.1 on page 5-42.*

5.2 Subroutines calls

A subroutine is a block of code that performs a task based on some arguments and optionally returns a result. By convention, registers R0 to R3 are used to pass arguments to subroutines, and R0 is used to pass a result back to the callers. A subroutine that needs more than 4 inputs uses the stack for the additional inputs.

To call subroutines, use a branch and link instruction. The syntax is:

```
BL destination
```

where *destination* is usually the label on the first instruction of the subroutine.

destination can also be a PC-relative expression.

The BL instruction:

- places the return address in the link register
- sets the PC to the address of the subroutine.

After the subroutine code is executed you can use a BX LR instruction to return.

———— **Note** ————

Calls between separately assembled or compiled modules must comply with the restrictions and conventions defined by the Procedure Call Standard for the ARM Architecture.

Example 5-1 shows a subroutine, doadd, that adds the values of two arguments and returns a result in R0.

Example 5-1 Add two arguments

```
AREA subrout, CODE, READONLY      ; Name this block of code
ENTRY                             ; Mark first instruction to execute
start MOV r0, #10                  ; Set up parameters
      MOV r1, #3
      BL doadd                    ; Call subroutine
stop  MOV r0, #0x18                ; angel_SWIreason_ReportException
      LDR r1, =0x20026             ; ADP_Stopped_ApplicationExit
      SVC #0x123456               ; ARM semihosting (formerly SWI)
doadd ADD r0, r0, r1              ; Subroutine code
      BX lr                      ; Return from subroutine
      END                       ; Mark end of file
```

5.2.1 See also

Concepts

- *Stack operations for nested subroutines* on page 5-30.

Reference

- *B, BL, BX, BLX, and BXJ* on page 3-130
- *Procedure Call Standard for the ARM Architecture* specification,
<http://infocenter/help/topic/com.arm.doc.ih0042->

5.3 Load immediates into registers

ARM and Thumb instructions can only be 32-bits wide. You can use the `MOV` and `MVN` instructions to load a register with an immediate value supported by the instruction set. Certain 32-bit values cannot be represented as an immediate operand to a single 32-bit instruction. These values can be loaded from memory in a single instruction.

In ARMv6T2 and later, you can load any 32-bit immediate value into a register with two instructions, a `MOV` followed by a `MOVT`. You can use a pseudo-instruction, `MOV32`, to construct the instruction sequence for you.

You can also use the `LDR` pseudo-instruction to load immediate values into a register.

You can include many commonly-used immediate values directly as operands within data processing instructions, without a separate load operation. The range of immediate values that you can include as operands in 16-bit Thumb instructions is much smaller.

5.3.1 See also

Concepts

- *Load immediate values using `MOV` and `MVN` on page 5-7*
- *Load 32-bit values to a register using `MOV32` on page 5-11*
- *Load immediate 32-bit values to a register using `LDR Rd, =const` on page 5-12*
- *Load values to VFP and NEON registers on page 9-10.*

5.4 Load immediate values using MOV and MVN

In ARM state:

- MOV can load any 8-bit immediate value, giving a range of 0x0-0xFF (0-255). It can also rotate these values by any even number. These values are also available as immediate operands in many data processing operations, without being loaded in a separate instruction.
- MVN can load the bitwise complements of these values. The numerical values are $-(n+1)$, where n is the value available in MOV.
- In ARMv6T2 and later, MOV can load any 16-bit number, giving a range of 0x0-0xFFFF (0-65535).

Table 5-1 shows the range of 8-bit values that can be loaded in a single ARM MOV or MVN instruction (for data processing operations). The value to load must be a multiple of the value shown in the Step column.

Table 5-2 on page 5-8 shows the range of 16-bit values that can be loaded in a single MOV ARM instruction in ARMv6T2 and later.

Table 5-1 ARM state immediate values (8-bit)

Binary	Decimal	Step	Hexadecimal	MVN value ^a	Notes
000000000000000000000000abcde fgh	0-255	1	0-0xFF	-1 to -256	-
000000000000000000000000abcde fgh00	0-1020	4	0-0x3FC	-4 to -1024	-
000000000000000000000000abcde fgh0000	0-4080	16	0-0xFF0	-16 to -4096	-
000000000000000000000000abcde fgh000000	0-16320	64	0-0x3FC0	-64 to -16384	-
...	-
abcde fgh000000000000000000000000000000	0-255 x 2 ²⁴	2 ²⁴	0-0xFFFF0000	1-256 x -2 ²⁴	-
cde fgh00000000000000000000000000ab	(bit pattern)	-	-	(bit pattern)	See b in Note
e f gh000000000000000000000000abcd	(bit pattern)	-	-	(bit pattern)	See b in Note
gh000000000000000000000000abcdef	(bit pattern)	-	-	(bit pattern)	See b in Note

Table 5-2 ARM state immediate values in MOV instructions

Binary	Decimal	Step	Hexadecimal	MVN value	Notes
0000000000000000abcdefghijklmnop	0-65535	1	0-0xFFFF	-	See c in Note

———— **Note** ————

These notes give extra information on Table 5-1 on page 5-7 and Table 5-2.

- | | |
|----------|--|
| a | The MVN values are only available directly as operands in MVN instructions. |
| b | These values are available in ARM state only. All the other values in this table are also available in 32-bit Thumb. |
| c | These values are only available in ARMv6T2 and later. They are not available directly as operands in other instructions. |

In Thumb state in ARMv6T2 and later:

- the 32-bit MOV instruction can load:
 - any 8-bit immediate value, giving a range of 0x0-0xFF (0-255)
 - any 8-bit immediate value, shifted left by any number
 - any 8-bit bit pattern duplicated in all four bytes of a register
 - any 8-bit bit pattern duplicated in bytes 0 and 2, with bytes 1 and 3 set to 0
 - any 8-bit bit pattern duplicated in bytes 1 and 3, with bytes 0 and 2 set to 0.These values are also available as immediate operands in many data processing operations, without being loaded in a separate instruction.
- the 32-bit MVN instruction can load the bitwise complements of these values. The numerical values are $-(n+1)$, where n is the value available in MOV.
- the 32-bit MOV instruction can load any 16-bit number, giving a range of 0x0-0xFFFF (0-65535). These values are not available as immediate operands in data processing operations.

In architectures with Thumb, the 16-bit Thumb MOV instruction can load any immediate value in the range 0-255.

Table 5-3 on page 5-9 shows the range of values that can be loaded in a single 32-bit Thumb MOV or MVN instruction (for data processing operations). The value to load must be a multiple of the value shown in the Step column.

Copyright © 2010 ARM. All rights reserved.
Non-Confidential

5-9

Binary	Decimal	Step	Hexadecimal	MVN value	Notes
0000000000000000abcdefghijklmnop	0-65535	1	0-0xFFFF	-	See c in Note

These notes give extra information on Table 5-3 and Table 5-4.

- In both ARM and Thumb, you do not have to decide whether to use MOV or MVN. The assembler uses whichever is appropriate. This is useful if the value is an assembly-time variable.

If you write an instruction with an immediate value that is not available, the assembler reports the error: Immediate *n* out of range for this operation.

5.4.1 See also

Concepts

- *Load immediates into registers* on page 5-6.

5.5 Load 32-bit values to a register using MOV32

In ARMv6T2 and later, both ARM and Thumb instruction sets include:

- a MOV instruction that can load any value in the range 0x00000000 to 0x0000FFFF into a register
- a MOVT instruction that can load any value in the range 0x0000 to 0xFFFF into the most significant half of a register, without altering the contents of the least significant half.

You can use these two instructions to construct any 32-bit immediate value in a register. Alternatively, you can use the MOV32 pseudo-instruction. The assembler generates the MOV, MOVT instruction pair for you.

You can also use the MOV32 instruction to load addresses into registers by using a label or any PC-relative expression in place of an immediate value. The assembler puts a relocation directive into the object file for the linker to resolve the address at link-time.

5.5.1 See also

Concepts

- *Register-relative and PC-relative expressions* on page 8-9.

Reference

Assembler Reference:

- *MOV32 pseudo-instruction* on page 3-180.

5.6 Load immediate 32-bit values to a register using LDR Rd, =const

The LDR Rd, =const pseudo-instruction can construct any 32-bit numeric value in a single instruction. You can use this pseudo-instruction to generate constants that are out of range of the MOV and MVN instructions.

The LDR pseudo-instruction generates the most efficient single instruction for the specified immediate value:

- If the immediate value can be constructed with a single MOV or MVN instruction, the assembler generates the appropriate instruction.
- If the immediate value cannot be constructed with a single MOV or MVN instruction, the assembler:
 - places the value in a *literal pool* (a portion of memory embedded in the code to hold constant values)
 - generates an LDR instruction with a PC-relative address that reads the constant from the literal pool.

For example:

```
LDR      rn, [pc, #offset to literal pool]
                                ; load register n with one word
                                ; from the address [pc + offset]
```

You must ensure that there is a literal pool within range of the LDR instruction generated by the assembler.

5.6.1 See also

Concepts

- *Literal pools* on page 5-13.

Reference

Assembler Reference:

- *LDR pseudo-instruction* on page 3-182.

5.7 Literal pools

The assembler uses literal pools to hold certain constant values that are to be loaded into registers. The assembler places a literal pool at the end of each section. The end of a section is defined either by the `END` directive at the end of the assembly or by the `AREA` directive at the start of the following section. The `END` directive at the end of an included file does not signal the end of a section.

In large sections the default literal pool can be out of range of one or more LDR instructions. The offset from the PC to the constant must be:

- less than 4KB in ARM or 32-bit Thumb code, but can be in either direction
- forward and less than 1KB in 16-bit Thumb LDR instruction.

When an `LDR Rd,=const` pseudo-instruction requires the immediate value to be placed in a literal pool, the assembler:

- Checks if the value is available and addressable in any previous literal pools. If so, it addresses the existing constant.
- Attempts to place the value in the next literal pool if it is not already available.

If the next literal pool is out of range, the assembler generates an error message. In this case you must use the `LTORG` directive to place an additional literal pool in the code. Place the `LTORG` directive after the failed LDR pseudo-instruction, and within $\pm 4\text{KB}$ (ARM, 32-bit Thumb-2) or in the range 0 to +1KB (16-bit Thumb).

You must place literal pools where the processor does not attempt to execute them as instructions. Place them after unconditional branch instructions, or after the return instruction at the end of a subroutine. Example 5-2 shows how this works.

The instructions listed as comments are the ARM instructions generated by the assembler.

Example 5-2 Placing literal pools

	AREA	Loadcon, CODE, READONLY	
	ENTRY		; Mark first instruction to execute
start	BL	func1	; Branch to first subroutine
	BL	func2	; Branch to second subroutine
stop	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
func1			

```

LDR    r0, =42                ; => MOV R0, #42
LDR    r1, =0x55555555        ; => LDR R1, [PC, #offset to
                                ; Literal Pool 1]
LDR    r2, =0xFFFFFFFF        ; => MVN R2, #0
BX     lr
LTORG                                ; Literal Pool 1 contains
                                ; literal 0x55555555

func2
LDR    r3, =0x55555555        ; => LDR R3, [PC, #offset to
                                ; Literal Pool 1]
; LDR r4, =0x66666666        ; If this is uncommented it
                                ; fails, because Literal Pool 2
                                ; is out of reach

BX     lr
LargeTable
SPACE  4200                    ; Starting at the current location,
                                ; clears a 4200 byte area of memory
                                ; to zero
END                                ; Literal Pool 2 is empty

```

5.7.1 See also

Concepts

- *Load immediate 32-bit values to a register using LDR Rd, =const* on page 5-12
- *LTORG* on page 6-20.

5.8 Load addresses into registers

It is often necessary to load an address into a register. You might have to load the address of a variable, a string literal, or the start location of a jump table.

Addresses are normally expressed as offsets from a label, or from the current PC or other register.

You can load an address into a register either:

- using the instruction `ADR`
- using the instruction `ADRL`
- using the instruction `MOV32`
- from a literal pool using the pseudo-instruction `LDR Rd, =Label`.

5.8.1 See also

Concepts

- *Load addresses to a register using `ADR` on page 5-16*
- *Load addresses to a register using `ADRL` on page 5-19*
- *Load 32-bit values to a register using `MOV32` on page 5-11*
- *Load addresses to a register using `LDR Rd, =label` on page 5-20.*

5.9 Load addresses to a register using ADR

The ADR instruction enables you to generate an address, within a certain range, without performing a data load. ADR accepts a PC-relative expression, that is, a label with an optional offset where the address of the label is relative to the current PC.

———— **Note** ————

The label used with ADR must be within the same code section. The assembler faults references to labels that are out of range in the same section.

The available range of addresses for the ADR instruction depends on the instruction set:

ARM	± 255 bytes to a byte or halfword-aligned address. ± 1020 bytes to a word-aligned address.
32-bit Thumb	± 4095 bytes to a byte, halfword, or word-aligned address.
16-bit Thumb	0 to 1020 bytes. <i>label</i> must be word-aligned. You can use the ALIGN directive to ensure this.

5.9.1 Example of a jump table implementation with ADR

Example 5-3 shows ARM code that implements a jump table. Here, the ADR instruction loads the address of the jump table.

Example 5-3 Implementing a jump table (ARM)

	AREA	Jump, CODE, READONLY	; Name this block of code
	ARM		; Following code is ARM code
num	EQU	2	; Number of entries in jump table
	ENTRY		; Mark first instruction to execute
start			; First instruction to call
	MOV	r0, #0	; Set up the three parameters
	MOV	r1, #3	
	MOV	r2, #2	
	BL	arithfunc	; Call the function
stop			
	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
arithfunc			; Label the function
	CMP	r0, #num	; Treat function code as unsigned
integer			
	BXHS	lr	; If code is >= num then simply return
	ADR	r3, JumpTable	; Load address of jump table


```

        LDR    pc, [r3,r0,LSL#2]        ; Jump to the appropriate routine
JumpTable
        DCD    DoAdd
        DCD    DoSub
DoAdd
        ADD    r0, r1, r2                ; Operation 0
        BX     lr                        ; Return
DoSub
        SUB    r0, r1, r2                ; Operation 1
        BX     lr                        ; Return
        END                                ; Mark the end of this file

```

In Example 5-3 on page 5-16, the function `arithfunc` takes three arguments and returns a result in `R0`. The first argument determines the operation to be carried out on the second and third arguments:

argument1=0 Result = argument2 + argument3.

argument1=1 Result = argument2 – argument3.

The jump table is implemented with the following instructions and assembler directives:

EQU	Is an assembler directive. It is used to give a value to a symbol. In Example 5-3 on page 5-16 it assigns the value 2 to <i>num</i> . When <i>num</i> is used elsewhere in the code, the value 2 is substituted. Using EQU in this way is similar to using <code>#define</code> to define a constant in C.
DCD	Declares one or more words of store. In Example 5-3 on page 5-16 each DCD stores the address of a routine that handles a particular clause of the jump table.
LDR	<p>The <code>LDR PC, [R3,R0,LSL#2]</code> instruction loads the address of the required clause of the jump table into the PC. It:</p> <ul style="list-style-type: none"> • multiplies the clause number in <code>R0</code> by 4 to give a word offset • adds the result to the address of the jump table • loads the contents of the combined address into the PC.

5.9.2 See also

Concepts

- *Load addresses to a register using LDR Rd, =label* on page 5-20
- *Load addresses to a register using ADR* on page 5-16.

Reference

Assembler Reference:

- *ADR (PC-relative)* on page 3-27.

5.10 Load addresses to a register using ADRL

The ADRL pseudo-instruction enables you to generate an address, within a certain range, without performing a data load. ADRL accepts a PC-relative expression, that is, a label with an optional offset where the address of the label is relative to the current PC.

Note

The label used with ADRL must be within the same code section. The assembler faults references to labels that are out of range in the same section.

ADRL is not available in Thumb state on processors before ARMv6T2.

The assembler converts an ADRL *rn, label* pseudo-instruction by generating:

- two data processing instructions that load the address, if it is in range
- an error message if the address cannot be constructed in two instructions.

The available range depends on the instruction set in use:

ARM	±64KB to a byte or halfword-aligned address. ±256KB to a word-aligned address.
32-bit Thumb	±1MB to a byte, halfword, or word-aligned address.
16-bit Thumb	ADRL is not available.

5.10.1 See also

Concepts

- *Load addresses to a register using LDR Rd, =label* on page 5-20
- *Load addresses to a register using ADR* on page 5-16.

5.11 Load addresses to a register using LDR Rd, =label

The LDR Rd,= pseudo-instruction can load any 32-bit numeric value into a register. It also accepts PC-relative expressions such as labels, and labels with offsets.

The assembler converts an LDR R0, =label pseudo-instruction by:

- placing the address of label in a literal pool (a portion of memory embedded in the code to hold constant values)
- generating a PC-relative LDR instruction that reads the address from the literal pool, for example:

```
LDR    rn [pc, #offset to literal pool]
                        ; load register n with one word
                        ; from the address [pc + offset]
```

You must ensure that there is a literal pool within range (see *Literal pools* on page 5-13 for more information).

Unlike the ADR and ADRL pseudo-instructions, you can use LDR with labels that are outside the current section. The assembler places a relocation directive in the object code when the source file is assembled. The relocation directive instructs the linker to resolve the address at link time. The address remains valid wherever the linker places the section containing the LDR and the literal pool.

Example 5-4 shows how this works.

The instructions listed in the comments are the ARM instructions generated by the assembler.

Example 5-4 Loading using LDR Rd, =label

	AREA	LDRlabel, CODE,READONLY	
	ENTRY		; Mark first instruction to execute
start	BL	func1	; Branch to first subroutine
	BL	func2	; Branch to second subroutine
stop	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
func1	LDR	r0, =start	; => LDR R0,[PC, #offset into ; Literal Pool 1]
	LDR	r1, =Darea + 12	; => LDR R1,[PC, #offset into ; Literal Pool 1]
	LDR	r2, =Darea + 6000	; => LDR R2, [PC, #offset into

```

                                ; Literal Pool 1]
                                ; Return
                                ; Literal Pool 1
func2    BX      1r
        LTORG
        LDR      r3, =Darea + 6000
                                ; => LDR r3, [PC, #offset into
                                ; Literal Pool 1]
                                ; (sharing with previous literal)
                                ; If uncommented produces an error
                                ; as Literal Pool 2 is out of range
        ; LDR      r4, =Darea + 6004
                                ; Return
Darea    BX      1r
        SPACE    8000
                                ; Starting at the current location,
                                ; clears a 8000 byte area of memory
                                ; to zero
        END
                                ; Literal Pool 2 is out of range of
                                ; the LDR instructions above

```

5.11.1 An LDR Rd, =label example: string copying

Example 5-5 shows an ARM code routine that overwrites one string with another string. It uses the LDR pseudo-instruction to load the addresses of the two strings from a data section. The following are particularly significant:

- DCB The DCB directive defines one or more bytes of store. In addition to integer values, DCB accepts quoted strings. Each character of the string is placed in a consecutive byte.
- LDR, STR The LDR and STR instructions use post-indexed addressing to update their address registers. For example, the instruction:

LDRB r2,[r1],#1

loads R2 with the contents of the address pointed to by R1 and then increments R1 by 1.

Example 5-5 String copy

	AREA	StrCopy, CODE, READONLY	
	ENTRY		; Mark first instruction to execute
start	LDR	r1, =srcstr	; Pointer to first string
	LDR	r0, =dststr	; Pointer to second string
	BL	strcpy	; Call subroutine to do copy
stop	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
strcpy	LDRB	r2, [r1],#1	; Load byte and update address
	STRB	r2, [r0],#1	; Store byte and update address
	CMP	r2, #0	; Check for zero terminator
	BNE	strcpy	; Keep going if not
	MOV	pc,lr	; Return
	AREA	Strings, DATA, READWRITE	
srcstr	DCB	"First string - source",0	
dststr	DCB	"Second string - destination",0	
	END		

5.11.2 See also

Concepts

- *Load immediate 32-bit values to a register using LDR Rd, =const* on page 5-12.

Reference

Assembler Reference:

- *LDR pseudo-instruction* on page 3-182
- *DCB* on page 6-25.

5.12 Other ways to Load and store registers

You can load any 32-bit value from memory into a register with an LDR data load instruction. To store registers into memory you can use the STR data store instruction.

You can use the MOV instruction to move any 32-bit data from one register to another.

5.12.1 See also

Concepts

- *Load and store multiple register instructions* on page 5-25
- *Load and store multiple instructions available in ARM and Thumb* on page 5-26.

Reference

Assembler Reference:

- *Memory access instructions* on page 3-12
- *MOV and MVN* on page 3-71.

5.13 Load and store multiple register instructions

The ARM and Thumb instruction sets include instructions that load and store multiple registers to and from memory.

Multiple register transfer instructions provide an efficient way of moving the contents of several registers to and from memory. They are most often used for block copy and for stack operations at subroutine entry and exit. The advantages of using a multiple register transfer instruction instead of a series of single data transfer instructions include:

- Smaller code size.
- A single instruction fetch overhead, rather than many instruction fetches.
- On uncached ARM processors, the first word of data transferred by a load or store multiple is always a nonsequential memory cycle, but all subsequent words transferred can be sequential memory cycles. Sequential memory cycles are faster in most systems.

Note

The lowest numbered register is transferred to or from the lowest memory address accessed, and the highest numbered register to or from the highest address accessed. The order of the registers in the register list in the instructions makes no difference.

You can use the `--diag_warning 1206` assembler command-line option to check that registers in register lists are specified in increasing order.

5.13.1 See also

Concepts

- *Load and store multiple instructions available in ARM and Thumb* on page 5-26
- *Stack implementation using LDM and STM* on page 5-28
- *Stack operations for nested subroutines* on page 5-30
- *Block copy with LDM and STM* on page 5-31.

5.14 Load and store multiple instructions available in ARM and Thumb

The following instructions are available in both ARM and Thumb instruction sets:

LDM	Load Multiple registers.
STM	Store Multiple registers.
PUSH	Store multiple registers onto the stack and update the stack pointer.
POP	Load multiple registers off the stack, and update the stack pointer.

In LDM and STM instructions:

- The list of registers loaded or stored can include:
 - in ARM instructions, any or all of R0-R12, SP, LR, and PC
 - in 32-bit Thumb instructions, any or all of R0-R12, and optionally LR or PC with some restrictions
 - in 16-bit Thumb instructions, any or all of R0-R7.
- The address can be:
 - incremented after each transfer
 - incremented before each transfer (ARM instructions only)
 - decremented after each transfer (ARM instructions only)
 - decremented before each transfer (not in 16-bit Thumb).
- The base register can be either:
 - updated to point to the next block of data in memory
 - left as it was before the instruction.

When the base register is updated to point to the next block in memory, this is called *writeback*, that is, the adjusted address is written back to the base register.

In PUSH and POP instructions:

- The stack pointer (SP) is the base register, and is always updated.
- The address is incremented after each transfer in POP instructions, and decremented before each transfer in PUSH instructions.
- The list of registers loaded or stored can include:
 - in ARM instructions, any or all of R0-R12, SP, LR, and PC
 - in 32-bit Thumb instructions, any or all of R0-R12, and optionally LR or PC with some restrictions
 - in 16-bit Thumb instructions, any or all of R0-R7, and optionally LR (PUSH only) or PC (POP only).

5.14.1 See also

Concepts

- *Load and store multiple register instructions* on page 5-25.

5.15 Stack implementation using LDM and STM

The load and store multiple instructions can update the base register. For stack operations, the base register is usually the stack pointer, SP. This means that you can use these instructions to implement push and pop operations for any number of registers in a single instruction.

The load and store multiple instructions can be used with several types of stack:

Descending or ascending

The stack grows downwards, starting with a high address and progressing to a lower one (a *descending* stack), or upwards, starting from a low address and progressing to a higher address (an *ascending* stack).

Full or empty

The stack pointer can either point to the last item in the stack (a *full* stack), or the next free space on the stack (an *empty* stack).

To make it easier for the programmer, stack-oriented suffixes can be used instead of the increment or decrement, and before or after suffixes. Table 5-5 shows the stack-oriented suffixes and their equivalent addressing mode suffixes for load and store instructions.

Table 5-5 Stack-oriented suffixes and equivalent addressing mode suffixes

Stack-oriented suffix	For store or push instructions	For load or pop instructions
FD (Full Descending stack)	DB (Decrement Before)	IA (Increment After)
FA (Full Ascending stack)	IB (Increment Before)	DA (Decrement After)
ED (Empty Descending stack)	DA (Decrement After)	IB (Increment Before)
EA (Empty Ascending stack)	IA (Increment After)	DB (Decrement Before)

Table 5-6 shows the load and store multiple instructions with the stack-oriented suffixes for the various stack types.

Table 5-6 Suffixes for load and store multiple instructions

Stack type	Store	Load
Full descending	STMFD (STMDB, Decrement Before)	LDMFD (LDM, increment after)

Table 5-6 Suffixes for load and store multiple instructions (continued)

Stack type	Store	Load
Full ascending	STMFA (STMIB, Increment Before)	LDMFA (LMDMA, Decrement After)
Empty descending	STMED (STMDA, Decrement After)	LDMED (LDMIB, Increment Before)
Empty ascending	STMEA (STM, increment after)	LDMEA (LDMDB, Decrement Before)

For example:

```

STMFD    sp!, {r0-r5} ; Push onto a Full Descending Stack
LDMFD    sp!, {r0-r5} ; Pop from a Full Descending Stack

```

Note

The *Procedure Call Standard for the ARM Architecture* (AAPCS), and ARM and Thumb C and C++ compilers always use a full descending stack.

The PUSH and POP instructions assume a full descending stack. They are the preferred synonyms for STMDB and LDM with writeback.

5.15.1 See also

Concepts

- *Load and store multiple register instructions* on page 5-25.

Reference

- *Procedure Call Standard for the ARM Architecture*,
<http://infocenter.arm.com/help/topic/com.arm.doc.ih0042-/index.html>.

5.16 Stack operations for nested subroutines

Stack operations are very useful at subroutine entry and exit. At the start of a subroutine, any working registers required can be stored on the stack, and at exit they can be popped off again.

In addition, if the link register is pushed onto the stack at entry, additional subroutine calls can be made safely without causing the return address to be lost. If you do this, you can also return from a subroutine by popping PC off the stack at exit, instead of popping LR and then moving that value into PC. For example:

```
subroutine  PUSH    {r5-r7,lr} ; Push work registers and lr
           ; code
           BL      somewhere_else
           ; code
           POP     {r5-r7,pc} ; Pop work registers and pc
```

———— Note ————

Use this with care in mixed ARM and Thumb systems. In ARMv4T systems, you cannot change state by popping directly into PC. In these cases you must pop the address into a temporary register and use the BX instruction.

In ARMv5T and later, you can change state in this way.

5.16.1 See also

Concepts

- *Subroutines calls* on page 5-4.

Developing Software for ARM processors:

- Chapter 5 *Interworking ARM and Thumb*.

5.17 Block copy with LDM and STM

Example 5-6 is an ARM code routine that copies a set of words from a source location to a destination by copying a single word at a time.

Example 5-6 Block copy without LDM and STM

	AREA	Word, CODE, READONLY	; name this block of code
num	EQU	20	; set number of words to be copied
	ENTRY		; mark the first instruction called
start			
	LDR	r0, =src	; r0 = pointer to source block
	LDR	r1, =dst	; r1 = pointer to destination block
	MOV	r2, #num	; r2 = number of words to copy
wordcopy			
	LDR	r3, [r0], #4	; load a word from the source and
	STR	r3, [r1], #4	; store it to the destination
	SUBS	r2, r2, #1	; decrement the counter
	BNE	wordcopy	; ... copy more
stop			
	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
	AREA	BlockData, DATA, READWRITE	
src	DCD	1,2,3,4,5,6,7,8,1,2,3,4,5,6,7,8,1,2,3,4	
dst	DCD	0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0	
	END		

This module can be made more efficient by using LDM and STM for as much of the copying as possible. Eight is a sensible number of words to transfer at a time, given the number of registers that the ARM has. The number of eight-word multiples in the block to be copied can be found (if R2 = number of words to be copied) using:

```
MOVS    r3, r2, LSR #3    ; number of eight word multiples
```

This value can be used to control the number of iterations through a loop that copies eight words per iteration. When there are less than eight words left, the number of words left can be found (assuming that R2 has not been corrupted) using:

```
ANDS    r2, r2, #7
```

Example 5-7 on page 5-32 lists the block copy module rewritten to use LDM and STM for copying.

Example 5-7 Block copy using LDM and STM

num	AREA	Block, CODE, READONLY	; name this block of code
	EQU	20	; set number of words to be copied
	ENTRY		; mark the first instruction called
start			
	LDR	r0, =src	; r0 = pointer to source block
	LDR	r1, =dst	; r1 = pointer to destination block
	MOV	r2, #num	; r2 = number of words to copy
blockcopy	MOV	sp, #0x400	; Set up stack pointer (sp)
	MOVS	r3,r2, LSR #3	; Number of eight word multiples
	BEQ	copywords	; Less than eight words to move?
	PUSH	{r4-r11}	; Save some working registers
octcopy	LDM	r0!, {r4-r11}	; Load 8 words from the source
	STM	r1!, {r4-r11}	; and put them at the destination
	SUBS	r3, r3, #1	; Decrement the counter
	BNE	octcopy	; ... copy more
	POP	{r4-r11}	; Don't need these now - restore ; originals
copywords	ANDS	r2, r2, #7	; Number of odd words to copy
	BEQ	stop	; No words left to copy?
wordcopy			
	LDR	r3, [r0], #4	; Load a word from the source and
	STR	r3, [r1], #4	; store it to the destination
	SUBS	r2, r2, #1	; Decrement the counter
stop	BNE	wordcopy	; ... copy more
	MOV	r0, #0x18	; angel_SWIreason_ReportException
	LDR	r1, =0x20026	; ADP_Stopped_ApplicationExit
	SVC	#0x123456	; ARM semihosting (formerly SWI)
	AREA	BlockData, DATA, READWRITE	
src	DCD	1,2,3,4,5,6,7,8,1,2,3,4,5,6,7,8,1,2,3,4	
dst	DCD	0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0	
	END		

5.18 Read-Modify-Write procedure

When you want to modify specific bits in a system register, you must ensure that you do not modify the other bits in the same register. This is because individual bits in a system register control different system functionality, and modifying them might cause your program to behave incorrectly. You must use a read-modify-write procedure to ensure that you modify only the bits you want to change.

To read-modify-write a system register, the instruction sequence must be:

1. The first instruction copies the value from the target system register to a temporary general-purpose register.
2. The next one or more instructions modify the required bits in the general-purpose register. This can be one or both of:
 - BIC to clear just the bits that need to be cleared to 0
 - ORR to set just the bits that need to be set to 1.
3. The final instruction writes the value from the general-purpose register to the target system register.

5.18.1 Example

This example shows the read-modify-write procedure to change some bits of a NEON and VFP system register FPSCR without affecting other bits

```
VMRS    r10,FPSCR           ; copy FPSCR into the general-purpose r10
BIC     r10,r10,#0x00370000 ; clears STRIDE bits[21:20] and LEN bits[18:16]
ORR     r10,r10,#0x00030000 ; sets bits[17:16] (STRIDE =1 and LEN = 4)
VMSR    FPSCR,r10           ; copy r10 back into FPSCR
```

5.18.2 See also

Reference

Assembler Reference:

- *VMRS and VMSR* on page 4-17.

5.19 Optional hash

You do not need to specify a hash before immediate constants in any instruction syntax (including ARM, Thumb, Wireless MMX, NEON, and VFP instructions). For example, the following are valid instructions:

```
BKPT 100
MOVT R1, 256
VCEQ.I8 Q1, Q2, 0
```

By default, the assembler warns if you do not specify a hash:

WARNING: A1865W: '#' not seen before constant expression.

This can be suppressed with `--diag_suppress=1865`.

If you use the assembler code with another assembler, you are advised to use the # before all immediates. The disassembler will always show the # for clarity.

5.19.1 See also

Reference

Assembler Reference:

- *Instruction summary* on page 3-3
- *NEON instructions* on page 4-2
- *Shared NEON and VFP instructions* on page 4-5
- *VFP instructions* on page 4-6.

5.20 Use of macros

A macro definition is a block of code enclosed between `MACRO` and `MEND` directives. It defines a name that can be used instead of repeating the whole block of code. The main uses for a macro are:

- to make it easier to follow the logic of the source code by replacing a block of code with a single meaningful name
- to avoid repeating a block of code several times.

5.20.1 See also

Concepts

- *Test-and-branch macro example* on page 5-36
- *Unsigned integer division macro example* on page 5-37.

Reference

Assembler Reference:

- *MACRO and MEND* on page 6-40.

5.21 Test-and-branch macro example

In ARM code in any processor and in Thumb code in processors before ARMv6T2, a test-and-branch operation requires two instructions to implement.

You can define a macro definition such as this:

```
MACRO
$label TestAndBranch $dest, $reg, $cc
$label CMP    $reg, #0
      B$cc    $dest
MEND
```

The line after the MACRO directive is the *macro prototype statement*. This defines the name (TestAndBranch) you use to invoke the macro. It also defines *parameters* (\$label, \$dest, \$reg, and \$cc). Unspecified parameters are substituted with an empty string. For this macro you must give values for \$dest, \$reg and \$cc to avoid syntax errors. The assembler substitutes the values you give into the code.

This macro can be invoked as follows:

```
test    TestAndBranch    NonZero, r0, NE
      ...
      ...
NonZero
```

After substitution this becomes:

```
test    CMP    r0, #0
      BNE    NonZero
      ...
      ...
NonZero
```

5.21.1 See also

Concepts

- *Use of macros* on page 5-35
- *Unsigned integer division macro example* on page 5-37
- *Local labels* on page 8-15.

5.22 Unsigned integer division macro example

Example 5-8 shows a macro that performs an unsigned integer division. It takes four parameters:

\$Bot	The register that holds the divisor.
\$Top	The register that holds the dividend before the instructions are executed. After the instructions are executed, it holds the remainder.
\$Div	The register where the quotient of the division is placed. It can be NULL ("") if only the remainder is required.
\$Temp	A temporary register used during the calculation.

Example 5-8 Unsigned integer division with a macro

```

$Lab    MACRO
DivMod  $Div,$Top,$Bot,$Temp
  ASSERT $Top <> $Bot          ; Produce an error message if the
  ASSERT $Top <> $Temp          ; registers supplied are
  ASSERT $Bot <> $Temp          ; not all different
  IF    "$Div" <> ""
    ASSERT $Div <> $Top        ; These three only matter if $Div
    ASSERT $Div <> $Bot        ; is not null ("")
    ASSERT $Div <> $Temp        ;
  ENDF
$Lab    MOV    $Temp, $Bot      ; Put divisor in $Temp
        CMP    $Temp, $Top, LSR #1 ; double it until
90      MOVLS  $Temp, $Temp, LSL #1 ; 2 * $Temp > $Top
        CMP    $Temp, $Top, LSR #1
        BLS    %b90           ; The b means search backwards
        IF    "$Div" <> ""      ; Omit next instruction if $Div is null
          MOV    $Div, #0       ; Initialize quotient
        ENDF
91      CMP    $Top, $Temp      ; Can we subtract $Temp?
        SUBCS  $Top, $Top,$Temp ; If we can, do so
        IF    "$Div" <> ""      ; Omit next instruction if $Div is null
          ADC    $Div, $Div, $Div ; Double $Div
        ENDF
        MOV    $Temp, $Temp, LSR #1 ; Halve $Temp,
        CMP    $Temp, $Bot      ; and loop until
        BHS    %b91           ; less than divisor
        MEND

```

The macro checks that no two parameters use the same register. It also optimizes the code produced if only the remainder is required.

To avoid multiple definitions of labels if DivMod is used more than once in the assembler source, the macro uses local labels (90, 91).

Example 5-9 shows the code that this macro produces if it is invoked as follows:

```
ratio DivMod R0,R5,R4,R2
```

Example 5-9 Output from division macro

	ASSERT	r5 <> r4	; Produce an error if the
	ASSERT	r5 <> r2	; registers supplied are
	ASSERT	r4 <> r2	; not all different
	ASSERT	r0 <> r5	; These three only matter if \$Div
	ASSERT	r0 <> r4	; is not null ("")
	ASSERT	r0 <> r2	;
ratio	MOV	r2, r4	; Put divisor in \$Temp
	CMP	r2, r5, LSR #1	; double it until
90	MOVLS	r2, r2, LSL #1	; 2 * r2 > r5
	CMP	r2, r5, LSR #1	
	BLS	%b90	; The b means search backwards
	MOV	r0, #0	; Initialize quotient
91	CMP	r5, r2	; Can we subtract r2?
	SUBCS	r5, r5, r2	; If we can, do so
	ADC	r0, r0, r0	; Double r0
	MOV	r2, r2, LSR #1	; Halve r2,
	CMP	r2, r4	; and loop until
	BHS	%b91	; less than divisor

5.22.1 See also

Concepts

- *Use of macros* on page 5-35
- *Test-and-branch macro example* on page 5-36
- *Local labels* on page 8-15.

5.23 Symbol versions

The ARM linker conforms to the *Base Platform ABI for the ARM Architecture* (BPABI) and supports the GNU-extended symbol versioning model.

To add a symbol version to an existing symbol, you must define a version symbol at the same address. A version symbol is of the form:

- `name@ver` if `ver` is a non default version of `name`
- `name@@ver` if `ver` is the default version of `name`.

The version symbols must be enclosed in vertical bars.

For example, to define a default version:

```
|my_versioned_symbol@@ver2|    ; Default version
my_asm_function PROC
    ...
    BX lr
ENDP
```

To define a non default version:

```
|my_versioned_symbol@ver1|    ; Non default version
my_old_asm_function PROC
    ...
    BX lr
ENDP
```

5.23.1 See also

Concepts

Using the Linker:

- Chapter 7 *Accessing image symbols*.

Reference

Base Platform ABI for the ARM Architecture,
<http://infocenter.arm.com/help/topic/com.arm.doc.ih0037-/index.html>.

5.24 Frame directives

You must use frame directives to describe the way that your code uses the stack if you want to be able to do either of the following:

- debug your application using stack unwinding
- use either flat or call-graph profiling.

The assembler uses frame directives to insert DWARF debug frame information into the object file in ELF format that it produces. This information is required by a debugger for stack unwinding and for profiling.

Be aware of the following:

- Frame directives do not affect the code produced by the assembler.
- The assembler does not validate the information in frame directives against the instructions emitted.

5.24.1 See also

Concepts

- *Exception tables and Unwind tables* on page 5-41.

Reference

- *Procedure Call Standard for the ARM Architecture*,
<http://infocenter.arm.com/help/topic/com.arm.doc.ih0042-/index.html>.

Assembler Reference:

- *Frame directives* on page 6-50.

5.25 Exception tables and Unwind tables

Exception tables are necessary to handle exceptions thrown by functions in high-level languages such as C++. *Unwind* tables contain debug frame information which are also necessary for the handling of such exceptions. An exception can only propagate through a function with an *unwind* table.

Functions written in C++ will have unwind information by default. However, for assembly language functions (code encased by PROC/ENDP or FUNC/ENDFUNC) that are called from C++ code, you must ensure that there are exception tables and *unwind* tables to enable the exceptions to propagate through them.

An exception cannot propagate through a function with a *nounwind* table. The exception handling runtime environment terminates the program if it encounters a *nounwind* table during exception processing.

The assembler can generate *nounwind* table entries for all functions and non-functions. The assembler can generate an *unwind* table for a function only if the function contains sufficient FRAME directives to describe the use of the stack within the function. To be able to create an *unwind* table for a function, each POP or PUSH instruction must be followed by a FRAME POP or FRAME PUSH directive respectively. Functions must conform to the conditions set out in the *Exception Handling ABI for the ARM Architecture* (EHABI), section 9.1 *Constraints on Use*. If the assembler cannot generate an *unwind* table it generates a *nounwind* table.

5.25.1 See also

Concepts

- *Frame directives* on page 5-40.

Reference

Assembler Reference:

- *Frame directives* on page 6-50
- *--no_exceptions_unwind* on page 2-23
- *--exceptions* on page 2-15
- *--no_exceptions* on page 2-22
- *FRAME UNWIND ON* on page 6-63
- *FRAME UNWIND OFF* on page 6-63
- *FUNCTION or PROC* on page 6-65
- *ENDFUNC or ENDP* on page 6-66.

5.26 Assembly language changes after RVCTv2.1

The assembly language accepted in RVCT v2.1 assembler and earlier is called pre-UAL ARM and Thumb. The current assembler accepts the UAL and the pre-UAL ARM and Thumb syntax. The assembler accepts the pre-UAL Thumb syntax only if it is preceded by a CODE16 directive, or if the source file is assembled with the --16 command-line option.

For the convenience of programmers who are familiar with the ARM and Thumb assembly languages accepted in RVCT v2.1 and earlier, Table 5-7 highlights the differences between the UAL and pre-UAL ARM assembly language syntax.

Table 5-7 Changes from earlier ARM assembly language

Change	Pre-UAL ARM syntax	Preferred UAL syntax
The default addressing mode for LDM and STM is IA	LDMIA, STMIA	LDM, STM
You can use the PUSH and POP mnemonics for full, descending stack operations in ARM as well as Thumb.	STMFD <i>sp!</i> , { <i>reglist</i> } LDMFD <i>sp!</i> , { <i>reglist</i> }	PUSH { <i>reglist</i> } POP { <i>reglist</i> }
You can use the LSL, LSR, ASR, ROR, and RRX instruction mnemonics for instructions with rotations and no other operation, in ARM as well as Thumb.	MOV <i>Rd</i> , <i>Rn</i> , LSL <i>shift</i> MOV <i>Rd</i> , <i>Rn</i> , LSR <i>shift</i> MOV <i>Rd</i> , <i>Rn</i> , ASR <i>shift</i> MOV <i>Rd</i> , <i>Rn</i> , ROR <i>shift</i> MOV <i>Rd</i> , <i>Rn</i> , RRX	LSL <i>Rd</i> , <i>Rn</i> , <i>shift</i> LSR <i>Rd</i> , <i>Rn</i> , <i>shift</i> ASR <i>Rd</i> , <i>Rn</i> , <i>shift</i> ROR <i>Rd</i> , <i>Rn</i> , <i>shift</i> RRX <i>Rd</i> , <i>Rn</i>
Use the <i>label</i> form for PC-relative addressing. Do not use the <i>offset</i> form in new code.	LDR <i>Rd</i> , [<i>pc</i> , # <i>offset</i>]	LDR <i>Rd</i> , <i>label</i>
Specify both registers for doubleword memory accesses. You must still obey rules about the register combinations you can use.	LDRD <i>Rd</i> , <i>addr_mode</i>	LDRD <i>Rd</i> , <i>Rd2</i> , <i>addr_mode</i>
{ <i>cond</i> }, if used, is always the last element of all instructions.	ADD{ <i>cond</i> }S LDR{ <i>cond</i> }SB	ADD{ <i>cond</i> } LDRSB{ <i>cond</i> }

In addition, some flexibility is permitted that was not permitted in previous assemblers as Table 5-8 shows.

Table 5-8 Relaxation of requirements

Relaxation	Permitted syntax	Preferred syntax
If the destination register is the same as the first operand, you can use a two register form of the instruction.	ADD <i>r1</i> , <i>r3</i>	ADD <i>r1</i> , <i>r1</i> , <i>r3</i>

You can write source code for Thumb processors earlier than ARMv6T2 using UAL.

If you are writing Thumb code for a processor earlier than ARMv6T2, you must restrict yourself to instructions that are available on the processor. The assembler generates error messages if you attempt to use an instruction that is not available.

If you are writing Thumb code for an ARMv6T2 or later processor, you can minimize your code size by using 16-bit instructions wherever possible.

Table 5-9 shows the main differences between the UAL and the pre-UAL Thumb assembly language.

Table 5-9 Differences between pre-UAL Thumb syntax and UAL syntax

Change	Pre-UAL Thumb syntax	UAL syntax
The default addressing mode for LDM and STM is IA	LDMIA, STMIA	LDM, STM
You must use the S postfix on instructions that update the flags. This change is essential to avoid conflict with 32-bit Thumb-2 instructions.	ADD r1, r2, r3 SUB r4, r5, #6 MOV r0, #1 LSR r1, r2, #1	ADDS r1, r2, r3 SUBS r4, r5, #6 MOVS r0, #1 LSRS r1, r2, #1
The preferred form for ALU instructions specifies three registers, even if the destination register is the same as the first operand. However, the UAL syntax allows the two register syntax.	ADD r7, r8 SUB r1, #80	ADD r7, r7, r8 SUBS r1, r1, #80
If <i>Rd</i> and <i>Rn</i> are both Lo registers, MOV <i>Rd</i> , <i>Rn</i> is disassembled as ADDS <i>Rd</i> , <i>Rn</i> , #0.	MOV r2, r3 MOV r8, r9 CPY r0, r1 LSL r2, r3, #0	ADDS r2, r3, #0 MOV r8, r9 MOV r0, r1 MOVS r2, r3
NEG <i>Rd</i> , <i>Rm</i> is disassembled as RSBS <i>Rd</i> , <i>Rm</i> , #0.	NEG <i>Rd</i> , <i>Rm</i>	RSBS <i>Rd</i> , <i>Rm</i> , #0

5.26.1 See also

Reference

Assembler Reference:

- *ARM, THUMB, THUMBX, CODE16 and CODE32* on page 6-75.

Chapter 6

Condition Codes

The following topics describe condition codes and conditional execution of ARM and Thumb code:

- *Conditional instructions* on page 6-2
- *Conditional execution in ARM state* on page 6-3
- *Conditional execution in Thumb state* on page 6-4
- *Updates to the ALU status flags* on page 6-5
- *Condition code suffixes* on page 6-7
- *Condition code meanings* on page 6-9
- *Benefits of using conditional execution* on page 6-11
- *Illustration of the benefits of using conditional instructions* on page 6-12
- *Optimization for execution speed* on page 6-16.

6.1 Conditional instructions

You can execute an instruction conditionally based upon the ALU status flags set by another instruction, either:

- immediately after the instruction that updated the flags
- after any number of intervening instructions that have not updated the flags.

The instructions that you can make conditional depends on whether the processor is in ARM state or Thumb state.

To make an instruction conditional, you must add a condition code suffix to the instruction mnemonic. The condition code suffix enables the processor to test a condition based on the flags. If the condition test of a conditional instruction fails, the instruction:

- does not execute
- does not write any value to its destination register
- does not affect any of the flags
- does not generate any exception.

6.1.1 See also

Concepts

- *Condition code suffixes* on page 6-7
- *Conditional execution in ARM state* on page 6-3
- *Conditional execution in Thumb state* on page 6-4.

6.2 Conditional execution in ARM state

Almost all ARM instructions can be executed conditionally on the value of the ALU status flags in the APSR. You can either add a condition code suffix to the instruction or you can conditionally skip over the instruction using a conditional branch instruction.

Using conditional branches instructions to control the flow of execution can be better when a series of instructions depend on the same condition.

Example 6-1 Conditional instructions to control execution

```
; flags set by a previous instruction
LSLEQ r0, r0, #24
ADDEQ r0, r0, #2
;...
```

Example 6-2 Conditional branch to control execution

```
; flags set by a previous instruction
BNE over
LSL r0, r0, #24
ADD r0, r0, #2
over
;...
```

6.2.1 See also

Concepts

- *Conditional execution in Thumb state* on page 6-4.

6.3 Conditional execution in Thumb state

In Thumb state on processors before ARMv6T2, the only mechanism for conditional execution is a conditional branch. You can conditionally skip over the instruction using a conditional branch instruction.

In Thumb state on ARMv6T2 or later processors, instructions can also be conditionally executed by:

- using CBZ and CBNZ
- using the IT (If-Then) instruction.

The Thumb CBZ (Conditional Branch on Zero) and CBNZ (Conditional Branch on Non-Zero) instructions compare the value of a register against zero and branch on the result.

IT is a 16-bit instruction that enables almost all Thumb instructions to be conditionally executed, on the value of the ALU flags, using the condition code suffix. Each IT instruction provides conditional execution for up to four following instructions.

Example 6-3 Conditional instructions using IT block

```

; flags set by a previous instruction
ITT    EQ
LSLEQ  r0, r0, #24
ADDEQ  r0, r0, #2
;...
```

6.3.1 See also

Concepts

- *Conditional execution in ARM state* on page 6-3.

Reference

- *IT* on page 3-134
- *CBZ and CBNZ* on page 3-137.

6.4 Updates to the ALU status flags

In ARM state, and in Thumb state on ARMv6T2 or later processors, most data processing instructions have an option to update ALU status flags in the *Application Program Status Register* (APSR) according to the result of the operation. Instructions with the optional S suffix update the condition flags. Conditional instructions that are not executed have no effect on the flags.

In Thumb state on processors before ARMv6T2, most data processing instructions update the ALU status flags automatically according to the result of the operation. There is no option to leave the flags unchanged and not update them. Other instructions cannot update the flags.

The APSR contains the following ALU status flags:

N	Set to 1 when the result of the operation is negative, cleared to 0 otherwise.
Z	Set to 1 when the result of the operation is zero, cleared to 0 otherwise.
C	Set to 1 when the operation results in a carry, cleared to 0 otherwise.
V	Set to 1 when the operation causes overflow, cleared to 0 otherwise.

A carry occurs:

- if the result of an addition is greater than or equal to 2^{32}
- if the result of a subtraction is positive or zero
- as the result of an inline barrel shifter operation in a move or logical instruction.

Overflow occurs if the result of an add, subtract, or compare is greater than or equal to 2^{31} , or less than -2^{31} .

The instruction also determines the flags that get updated. Some instructions update all flags, and some instructions only update a subset of the flags. If a flag is not updated, the original value is preserved. In the *Assembler Reference*, the description of each instruction details the effect it has on the flags.

Note

Most instructions update the status flags only if the S suffix is specified. The instructions CMP, CMN, TEQ, and TST always update the condition code flags.

6.4.1 See also

Concepts

- *Condition code suffixes* on page 6-7.

6.5 Condition code suffixes

The instructions that can be conditional have an optional condition code, shown in syntax descriptions as *{cond}*. This condition is encoded in ARM instructions, and encoded in a preceding IT instruction for Thumb instructions. An instruction with a condition code is only executed if the condition code flags in the APSR meet the specified condition.

In Thumb state on processors before ARMv6T2, the *{cond}* field is only permitted on certain branch instructions because there is no IT instruction on these processors.

The following table shows the condition codes that you can use and the flags they depend on.

Table 6-1 Condition code suffixes

Suffix	Flags	Meaning
EQ	Z set	Equal
NE	Z clear	Not equal
CS or HS	C set	Higher or same (unsigned \geq)
CC or LO	C clear	Lower (unsigned $<$)
MI	N set	Negative
PL	N clear	Positive or zero
VS	V set	Overflow
VC	V clear	No overflow
HI	C set and Z clear	Higher (unsigned $>$)
LS	C clear or Z set	Lower or same (unsigned \leq)
GE	N and V the same	Signed \geq
LT	N and V differ	Signed $<$
GT	Z clear, N and V the same	Signed $>$
LE	Z set, N and V differ	Signed \leq
AL	Any	Always. This suffix is normally omitted.

The following is an example of conditional execution.

Example 6-4

ADD	r0, r1, r2	; r0 = r1 + r2, don't update flags
ADDS	r0, r1, r2	; r0 = r1 + r2, and update flags
ADDSCS	r0, r1, r2	; If C flag set then r0 = r1 + r2, and update flags
CMP	r0, r1	; update flags based on r0-r1.

6.5.1 See also**Concepts**

- *Updates to the ALU status flags* on page 6-5
- *Condition code meanings* on page 6-9
- *Conditional instructions* on page 6-2.

6.6 Condition code meanings

The precise meanings of the condition code flags differ depending on whether the flags were set by a floating-point operation or by an ARM data processing instruction. This is because:

- floating-point values are never unsigned, so the unsigned conditions are not required
- *Not-a-Number* (NaN) values have no ordering relationship with numbers or with each other, so additional conditions are required to account for *unordered* results.

The only VFP instruction that can be used to update the status flags is `VCMP`. Other VFP or NEON instruction cannot modify the condition code flags.

The `VCMP` instruction does not update the flags in the APSR directly, but updates a separate set of flags in the FPSCR. To use these flags to control conditional instructions, including conditional VFP instructions, you must first copy them into the APSR using a `VMRS` instruction:

```
VMRS APSR_nzcv, FPSCR
```

The meanings of the condition code mnemonics are shown in Table 6-2.

Table 6-2 Condition codes

Mnemonic	Meaning after ARM data processing instruction	Meaning after VFP <code>VCMP</code> instruction
EQ	Equal	Equal
NE	Not equal	Not equal, or unordered
CS	Carry set	Greater than or equal, or unordered
HS	Unsigned higher or same	Greater than or equal, or unordered
CC	Carry clear	Less than
LO	Unsigned lower	Less than
MI	Negative	Less than
PL	Positive or zero	Greater than or equal, or unordered
VS	Overflow	Unordered (at least one NaN operand)
VC	No overflow	Not unordered
HI	Unsigned higher	Greater than, or unordered

Table 6-2 Condition codes (continued)

Mnemonic	Meaning after ARM data processing instruction	Meaning after VFP VCMP instruction
LS	Unsigned lower or same	Less than or equal
GE	Signed greater than or equal	Greater than or equal
LT	Signed less than	Less than, or unordered
GT	Signed greater than	Greater than
LE	Signed less than or equal	Less than or equal, or unordered
AL	Always (normally omitted)	Always (normally omitted)

———— **Note** ————

The type of the instruction that last updated the flags in the APSR determines the meaning of condition codes.

6.6.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *Condition code suffixes* on page 6-7.

Reference

Assembler Reference:

- *IT* on page 3-134
- *VMRS and VMSR* on page 4-17.

6.7 Benefits of using conditional execution

You can use conditional execution of ARM instructions to reduce the number of branch instructions in your code. This improves code density. The IT instruction in Thumb-2 achieves a similar improvement.

Branch instructions are also expensive in processor cycles. On ARM processors without branch prediction hardware, it typically takes three processor cycles to refill the processor pipeline each time a branch is taken.

Some ARM processors, for example ARM10™ and StrongARM®, have branch prediction hardware. In systems using these processors, the pipeline only has to be flushed and refilled when there is a misprediction.

6.7.1 See also

Concepts

- *Illustration of the benefits of using conditional instructions* on page 6-12.

6.8 Illustration of the benefits of using conditional instructions

This illustrates the difference between using branches and using conditional instructions. It uses the Euclid algorithm for the *Greatest Common Divisor* (gcd) to demonstrate how conditional instructions improve code size and speed.

In C the gcd algorithm can be expressed as:

```
int gcd(int a, int b)
{
    while (a != b)
    {
        if (a > b)
            a = a - b;
        else
            b = b - a;
    }
    return a;
}
```

The following examples show implementations of the gcd algorithm with and without conditional instructions.

Note

The detailed analysis of execution speed only applies to an ARM7™ processor. The code density calculations apply to all ARM processors.

6.8.1 Example of conditional execution using branches in ARM code

This is an ARM code implementation of the gcd algorithm using branches, without using any other conditional instructions. Conditional execution is achieved by using conditional branches, rather than individual conditional instructions:

```
gcd    CMP     r0, r1
      BEQ     end
      BLT     less
      SUBS    r0, r0, r1 ; could be SUB r0, r0, r1 for ARM
      B       gcd
less   SUBS    r1, r1, r0 ; could be SUB r1, r1, r0 for ARM
      B       gcd
end
```

The code is seven instructions long because of the number of branches. Every time a branch is taken, the processor must refill the pipeline and continue from the new location. The other instructions and non-executed branches use a single cycle each.

The following table shows the number of cycles this implementation uses on an ARM7 processor when R0 equals 1 and R1 equals 2.

Table 6-3 Conditional branches only

R0: a	R1: b	Instruction	Cycles (ARM7)
1	2	CMP r0, r1	1
1	2	BEQ end	1 (not executed)
1	2	BLT less	3
1	2	SUB r1, r1, r0	1
1	2	B gcd	3
1	1	CMP r0, r1	1
1	1	BEQ end	3
			Total = 13

6.8.2 Example of conditional execution using conditional instructions in ARM code

This is an ARM code implementation of the gcd algorithm using individual conditional instructions in ARM code. The gcd algorithm only takes four instructions:

```
gcd
    CMP    r0, r1
    SUBGT  r0, r0, r1
    SUBLE  r1, r1, r0
    BNE    gcd
```

In addition to improving code size, in most cases this code executes faster than the version that uses only branches.

The following table shows the number of cycles this implementation uses on an ARM7 processor when R0 equals 1 and R1 equals 2.

Table 6-4 All instructions conditional

R0: a	R1: b	Instruction	Cycles (ARM7)
1	2	CMP r0, r1	1
1	2	SUBGT r0,r0,r1	1 (not executed)
1	1	SUBLT r1,r1,r0	1

Table 6-4 All instructions conditional (continued)

R0: a	R1: b	Instruction	Cycles (ARM7)
1	1	BNE gcd	3
1	1	CMP r0,r1	1
1	1	SUBGT r0,r0,r1	1 (not executed)
1	1	SUBLT r1,r1,r0	1 (not executed)
1	1	BNE gcd	1 (not executed)
			Total = 10

Comparing this with the example that uses only branches:

- Replacing branches with conditional execution of all instructions saves three cycles.
- Where R0 equals R1, both implementations execute in the same number of cycles. For all other cases, the implementation that uses conditional instructions executes in fewer cycles than the implementation that uses branches only.

6.8.3 Example of conditional execution using conditional instructions in Thumb code

In architectures ARMv6T2 and later, you can use the IT instruction to write conditional instructions in Thumb code. The Thumb code implementation of the gcd algorithm using conditional instructions is very similar to the implementation in ARM code. The implementation in Thumb code is:

```
gcd
    CMP    r0, r1
    ITE    GT
    SUBGT  r0, r0, r1
    SUBLE  r1, r1, r0
    BNE    gcd
```

This assembles equally well to ARM or Thumb code. The assembler checks the IT instructions, but omits them on assembly to ARM code.

It requires one more instruction in Thumb code (the IT instruction) than in ARM code, but the overall code size is 10 bytes in Thumb code compared with 16 bytes in ARM code.

6.8.4 Example of conditional execution code using branches in Thumb code

In architectures before ARMv6T2, there is no IT instruction and hence Thumb instructions cannot be executed conditionally except for the B branch instruction. The gcd algorithm must be written with conditional branches and is very similar to the ARM code implementation using branches, without conditional instructions.

The Thumb code implementation of the gcd algorithm without conditional instructions requires seven instructions. The overall code size is 14 bytes. This is even less than the ARM implementation that uses conditional instructions, which uses 16 bytes.

In addition, on a system using 16-bit memory this Thumb implementation runs faster than both ARM implementations because only one memory access is required for each 16-bit Thumb instruction, whereas each 32-bit ARM instruction requires two fetches.

6.8.5 See also

Concepts

- *Benefits of using conditional execution* on page 6-11
- *Condition code suffixes* on page 6-7
- *Optimization for execution speed* on page 6-16.

Reference

ARM Architecture Reference Manual,
<http://infocenter.arm.com/help/topic/com.arm.doc.subset.arch.reference/index.html>.

Technical Reference Manual for your processor

Assembler Reference:

- *IT* on page 3-134.

6.9 Optimization for execution speed

To optimize code for execution speed you must have detailed knowledge of the instruction timings, branch prediction logic, and cache behavior of your target system.

6.9.1 See also

Reference

ARM Architecture Reference Manual,

<http://infocenter.arm.com/help/topic/com.arm.doc.subset.arch.reference/index.html>.

Technical Reference Manual for your processor.

Chapter 7

Using the Assembler

The following topics describe how to use the Assembler:

- *Assembler command line syntax* on page 7-2
- *Assembler commands listed in groups* on page 7-3
- *Specify command-line options with an environment variable* on page 7-6
- *Built-in variables and constants* on page 7-7
- *Versions of armasm* on page 7-12
- *Diagnostic messages* on page 7-13
- *Interlocks diagnostics* on page 7-14
- *IT block generation* on page 7-15
- *Thumb branch target alignment* on page 7-16
- *Instruction width* on page 7-17
- *Using the C preprocessor* on page 7-18
- *Address alignment* on page 7-20
- *Instruction width selection in Thumb* on page 7-21.

7.1 Assembler command line syntax

The command for invoking the assembler is:

```
armasm {options} {inputfile}
```

where *inputfile* is an assembly source file and *options* instruct the assembler how to assemble the *inputfile*. You can invoke the assembler with any combination options separated by spaces.

Note

The inline and embedded assemblers are part of the C and C++ compilers and do not use any command line syntax for invocation. However, to pass additional assembler options when the compiler invokes `armasm` for embedded assembly, you can use the `armcc -A` option.

The assembler command line is case-insensitive, except in filenames and where specified. The assembler uses the normal command line ordering rules as described in the *Using the Compiler*. Therefore, if the command line contains options that conflict with each other, then the last option found always takes precedence.

7.1.1 See also

Concepts

Using the Compiler:

- *Order of compiler command-line options* on page 3-12
- *Compiler command-line options listed by group* on page 3-4.

Reference

- *Assembler commands listed in groups* on page 7-3.

7.2 Assembler commands listed in groups

See the following command-line options in the *Assembler Reference*.

Help

- `--help` on page 2-19
- `--version_number` on page 2-29
- `--vsn` on page 2-30.

Source

- `--16` on page 2-5
- `--32` on page 2-5
- `--arm` on page 2-7
- `--arm_only` on page 2-7
- `-idir{,dir, ...}` on page 2-19
- `--maxcache=n` on page 2-22
- `--no_esc` on page 2-22
- `--no_regs` on page 2-23
- `--pd` on page 2-25
- `--predefine "directive"` on page 2-25
- `--reduce_paths` on page 2-26
- `--regnames=none` on page 2-26
- `--regnames=callstd` on page 2-26
- `--regnames=all` on page 2-27
- `--thumb` on page 2-28
- `--unsafe` on page 2-29.

Output

- `--debug` on page 2-10
- `--depend=dependfile` on page 2-11
- `--depend_format=string` on page 2-11
- `--dllexport_all` on page 2-14
- `--dwarf2` on page 2-15
- `--dwarf3` on page 2-15
- `-g` on page 2-19
- `--keep` on page 2-19

- *--length=n* on page 2-19
- *--list=file* on page 2-21
- *-m* on page 2-22
- *--md* on page 2-22
- *--no_code_gen* on page 2-22
- *--no_hide_all* on page 2-23
- *--no_terse* on page 2-24
- *-o filename* on page 2-24
- *--width=n* on page 2-30
- *--xref* on page 2-31
- *--untyped_local_labels* on page 2-29.

Target

- *--apcs=qualifier...qualifier* on page 2-5
- *--bi* on page 2-7
- *--bigend* on page 2-7
- *--compatible=name* on page 2-8
- *--cpu=name* on page 2-10
- *--device=list* on page 2-11
- *--device=name* on page 2-12
- *--fpmode=model* on page 2-16
- *--fpu=name* on page 2-17
- *--li* on page 2-20
- *--library_type=lib* on page 2-20
- *--littleend* on page 2-21
- *--no_unaligned_access* on page 2-24
- *--unaligned_access* on page 2-29.

Diagnostics

- *--brief_diagnostics* on page 2-8
- *--checkreglist* on page 2-8
- *--diag_error=tag{, tag}* on page 2-12
- *--diag_remark=tag{, tag}* on page 2-13
- *--diag_warning=tag{, tag}* on page 2-14

- `--diag_suppress=tag{, tag}` on page 2-13
- `--diag_style=style` on page 2-13
- `--errors=errorfile` on page 2-15
- `--no_warn` on page 2-24
- `--report-if-not-wysiwyg` on page 2-27
- `--split_ldm` on page 2-28.

Preprocessor

- `--cpreproc` on page 2-9
- `--cpreproc_opts=options` on page 2-9.

Exception table generation

- `--exceptions` on page 2-15
- `--exceptions_unwind` on page 2-15.

Project template

- `--project=filename` on page 2-25
- `--reinitialize_workdir` on page 2-27
- `--workdir=directory` on page 2-30.

License

- `--licretry` on page 2-20.

Command line options in a text file

- `--via=file` on page 2-30.

7.3 Specify command-line options with an environment variable

You can specify command-line options by setting the value of the ARMCC41_ASMOPT environment variable. The syntax is identical to the command line syntax. The assembler reads the value of ARMCC41_ASMOPT and inserts it at the front of the command string. This means that options specified in ARMCC41_ASMOPT can be overridden by arguments on the command-line.

7.3.1 See also

Concepts

- *Assembler command line syntax* on page 7-2.

7.4 Built-in variables and constants

Table 7-1 lists the built-in variables defined by the assembler.

Table 7-1 Built-in variables

{ARCHITECTURE}	Holds the name of the selected ARM architecture.
{AREANAME}	Holds the name of the current AREA.
{ARMASM_VERSION}	Holds an integer that increases with each version of armasm. The format of the version number is <i>PVbbbb</i> where: P is the major version V is the minor version bbbb is the build number.
ads\$version	Has the same value as {ARMASM_VERSION}.
{CODESIZE}	Is a synonym for {CONFIG}.
{COMMANDLINE}	Holds the contents of the command line.
{CONFIG}	Has the value 32 if the assembler is assembling ARM code, or 16 if it is assembling Thumb code.
{CPU}	Holds the name of the selected CPU. The default is "ARM7TDMI". If an architecture was specified in the command line --cpu option, {CPU} holds the value "Generic ARM".
{ENDIAN}	Has the value "big" if the assembler is in big-endian mode, or "little" if it is in little-endian mode.
{FPIC}	Has the boolean value True if /fpic is set. The default is False.
{FPU}	Holds the name of the selected FPU. The default is "SoftVFP".
{INPUTFILE}	Holds the name of the current source file.
{INTER}	Has the boolean value True if /inter is set. The default is False.
{LINENUM}	Holds an integer indicating the line number in the current source file.
{LINENUMUP}	When used in a macro, holds an integer indicating the line number of the current macro. The value is same as {LINENUM} when used in a non-macro context.
{LINENUMUPPER}	When used in a macro, holds an integer indicating the line number of the top macro. The value is same as {LINENUM} when used in a non-macro context.
{OPT}	Value of the currently-set listing option. The OPT directive can be used to save the current listing option, force a change in it, or restore its original value.
{PC} or .	Address of current instruction.

Table 7-1 Built-in variables (continued)

{PCSTOREOFFSET}	Is the offset between the address of the STR PC, [...] or STM Rb, {..., PC} instruction and the value of PC stored out. This varies depending on the CPU or architecture specified.
{ROPI}	Has the boolean value True if /ropi is set. The default is False.
{RWPI}	Has the boolean value True if /rwp i is set. The default is False.
{VAR} or @	Current value of the storage area location counter.

Built-in variables cannot be set using the SETA, SETL, or SETS directives. They can be used in expressions or conditions, for example:

```
IF {ARCHITECTURE} = "4T"
```

The built-in variable |ads\$version| must be all in lowercase. The names of the other built-in variables can be in uppercase, lowercase, or mixed. For example:

```
IF {CpU} = "Generic ARM"
```

———— **Note** ————

All built-in string variables contain case-sensitive values. Relational operations on these built-in variables will not match with strings that contain an incorrect case. Use the command line options --cpu and --fpu to determine valid values for {CPU}, {ARCHITECTURE}, and {FPU}.

Table 7-2 lists the built-in Boolean constants defined by the assembler.

Table 7-2 Built-in Boolean constants

{FALSE}	Logical constant false.
{TRUE}	Logical constant true.

Table 7-3 lists the target CPU related built-in variables that are predefined by the assembler. Where the value field is empty, the symbol is a boolean value and the meaning column describes when its value is {TRUE}.

Table 7-3 Predefined macros

Name	Value	Meaning
{TARGET_ARCH_ARM}	<i>num</i>	The number of the ARM base architecture of the target CPU irrespective of whether the assembler is assembling for ARM or Thumb. For possible values of {TARGET_ARCH_ARM} in relation to the ARM architecture versions, see Table 7-4 on page 7-10.
{TARGET_ARCH_THUMB}	<i>num</i>	The number of the Thumb base architecture of the target CPU irrespective of whether the assembler is assembling for ARM or Thumb. The value is defined as zero if the target does not support Thumb. For possible values of {TARGET_ARCH_THUMB} in relation to the ARM architecture versions, see Table 7-4 on page 7-10.
{TARGET_ARCH_XX}	–	XX represents the target architecture and its value depends on the target CPU. For example, if you specify the assembler option --cpu=4T or --cpu=ARM7TDMI then {TARGET_ARCH_4T} is defined. Table 7-4 on page 7-10 shows the possible values for XX.
{TARGET_FEATURE_EXTENSION_REGISTER_COUNT}	<i>num</i>	The number of 64-bit extension registers available in NEON or VFP.
{TARGET_FEATURE_CLZ}	–	If target CPU supports the CLZ instruction (that is, ARMv5T and later except ARMv6-M).
{TARGET_FEATURE_DIVIDE}	–	If the target CPU supports the hardware divide instructions SDIV and UDIV in Thumb (that is, ARMv7-M or ARMv7-R).
{TARGET_FEATURE_DOUBLEWORD}	–	If the target CPU supports the LDRD and STRD instructions (that is, ARMv5TE and later except ARMv6-M).
{TARGET_FEATURE_DSPMUL}	–	If the DSP-enhanced multiplier (for example the SMLAx instruction) is available, for example ARMv5TE.
{TARGET_FEATURE_MULTIPLY}	–	If the target CPU supports the long multiply instructions SMULL, SMLAL, UMULL, and UMLAL (that is, all architectures except ARMv6-M).
{TARGET_FEATURE_MULTIPROCESSING}	–	If assembling for a target CPU with ARMv7 Multiprocessing Extensions.
{TARGET_FEATURE_NEON}	–	If the target CPU has NEON.
{TARGET_FEATURE_NEON_FP16}	–	If the target CPU has NEON with half-precision floating-point operations.

Table 7-3 Predefined macros (continued)

Name	Value	Meaning
{TARGET_FEATURE_NEON_FP32}	—	If the target CPU has NEON with single-precision floating-point operations.
{TARGET_FEATURE_NEON_INTEGER}	—	If the target CPU has NEON with integer operations.
{TARGET_FEATURE_UNALIGNED}	—	If the target CPU support for unaligned access (that is, ARMv6 and later except ARMv6-M).
{TARGET_FPU_SOFTVFP}		If assembling with the option <code>--fpu=softvfp</code> .
{TARGET_FPU_SOFTVFP_VFP}		If assembling for a target CPU with softvfp and a hardware vfp, for example <code>--fpu=softvfp+vfpv3</code> .
{TARGET_FPU_VFP}		If assembling for a target CPU with a hardware VFP, without using softvfp, for example <code>--fpu=vfpv3</code>
{TARGET_FPU_VFPV2}		If assembling for a target CPU with VFPv2.
{TARGET_FPU_VFPV3}		If assembling for a target CPU with VFPv3.
{TARGET_PROFILE_A}		If assembling for a Cortex-A profile CPU (that is, <code>--cpu=7-A</code> option).
{TARGET_PROFILE_M}		If assembling for a Cortex-M profile CPU (that is, <code>--cpu=6-M</code> , <code>--cpu=6S-M</code> , or <code>--cpu=7-M</code>):
{TARGET_PROFILE_R}		If assembling for a Cortex-R profile CPU (that is, <code>--cpu=7-R</code> option).

Table 7-4 shows the possible values for {TARGET_ARCH_THUMB} (see Table 7-3 on page 7-9), and how these values relate to versions of the ARM architecture.

Table 7-4 {TARGET_ARCH_ARM} in relation to {TARGET_ARCH_THUMB}

ARM architecture	{TARGET_ARCH_ARM}	{TARGET_ARCH_THUMB}	xx
v4	4	0	4
v4T	4	1	4T
v5T	5	2	5T
v5TE	5	2	5TE
v5TEJ	5	2	5TEJ
v6	6	3	6
v6K	6	3	6K

Table 7-4 {TARGET_ARCH_ARM} in relation to {TARGET_ARCH_THUMB} (continued)

ARM architecture	{TARGET_ARCH_ARM}	{TARGET_ARCH_THUMB}	xx
v6Z	6	3	6Z
v6T2	6	4	6T2
v6-M	0	3	6M
v6S-M	0	3	6SM
v7-A	7	4	7A
v7-R	7	4	7R
v7-M	0	4	7M

7.4.1 See also

Reference

Assembler Reference:

- *--cpu=name* on page 2-10
- *--fpu=name* on page 2-17
- *Versions of armasm* on page 7-12.

7.5 Versions of `armasm`

You can use the built-in variable `{ARMASM_VERSION}` to distinguish between versions of `armasm`. The format of the version number is *PVbbbb* where:

P is the major version
V is the minor version
bbbb is the build number

The assembler did not have the built-in variable `|ads$version|` before ADS and RVCT. If you have to build versions of your code using legacy development tools, you can test for the built-in variable `|ads$version|`. If this variable is not defined, then the assembler is part of a legacy development toolchain. Use code similar to the following:

```
IF :DEF: |ads$version|
    ; code for RealView or ADS
ELSE
    ; code for SDT (a legacy development toolchain)
ENDIF
```

7.5.1 See also

Concepts

- *Built-in variables and constants* on page 7-7.

7.6 Diagnostic messages

In addition to the default error, warning and remark messages, the assembler can provide more diagnostic messages. By default, these additional diagnostic messages are not displayed. However, you can enable these additional messages using command-line options.

7.6.1 See also

Concepts

- *Interlocks diagnostics* on page 7-14
- *IT block generation* on page 7-15
- *Thumb branch target alignment* on page 7-16.

Reference

Assembler Reference:

- *--diag_error=tag{, tag}* on page 2-12.

7.7 Interlocks diagnostics

You can get warning messages about possible interlocks in your code caused by the pipeline of the processor chosen by the `--cpu` option. To do this, use the following command-line option when invoking the assembler:

```
armasm --diag_warning 1563
```

Note

Where the `--cpu` option specifies a multi-issue processor such as Cortex-A8, the assembler warnings are unpredictable.

7.7.1 See also

Concepts

- *Diagnostic messages* on page 7-13
- *IT block generation* on page 7-15
- *Thumb branch target alignment* on page 7-16
- *Instruction width* on page 7-17.

Reference

Assembler Reference:

- `--diag_error=tag{, tag}` on page 2-12.

7.8 IT block generation

If you write:

```
AREA x, CODE
THUMB
MOVNE    r0, r1 ; (1)
NOP
IT        NE
MOVNE    r0, r1 ; (2)
END
```

the assembler generates an IT instruction before the first MOVNE instruction.

You can get warning messages about this automatic generation of IT blocks when assembling Thumb code. To do this, use the following command-line option when invoking the assembler:

```
armasm --diag_warning 1763
```

7.8.1 See also

Concepts

- *Diagnostic messages* on page 7-13.

Reference

Assembler Reference:

- `--diag_error=tag{, tag}` on page 2-12.

7.9 Thumb branch target alignment

On some processors, non word-aligned Thumb instructions sometimes take one or more additional cycles to execute in loops. This means that it can be an advantage to ensure that branch targets are word-aligned. The assembler can issue warnings when branch targets in Thumb code are not word-aligned. To do this, use the following command-line option when invoking the assembler:

```
armasm --diag_warning 1604
```

7.9.1 See also

Concepts

- *Diagnostic messages* on page 7-13.

Reference

Assembler Reference:

- `--diag_error=tag{, tag}` on page 2-12.

7.10 Instruction width

If you specify the `.w` specifier, the instruction is encoded in 32 bits even if it can be encoded in 16 bits. You can use a diagnostic warning to detect when a branch instruction could have been encoded in 16 bits, but has been encoded in 32 bits. To do this, use the following command-line option when invoking the assembler:

```
armasm --diag_warning 1607
```

Note

This diagnostic does not produce a warning for relocated branch instructions, because the final address is not known. The linker might even insert a veneer, if the branch is out of range for a 32-bit instruction.

7.10.1 See also

Concepts

- *Diagnostic messages* on page 7-13.

Reference

Assembler Reference:

- `--diag_error=tag{, tag}` on page 2-12.

7.11 Using the C preprocessor

You can use C preprocessor commands in your assembly language source file. If you do this, you must use the `--cpreproc` command-line option when invoking the assembler. This causes `armasm` to call `armcc` to preprocess the file before assembling it.

`armasm` looks for the `armcc` binary in the same directory as the `armasm` binary. If it does not find the binary, it expects it to be on the `PATH`.

`armasm` passes certain options to `armcc` if present on the command-line. These are shown in Table 7-5. Some of these options are converted to the `armcc` equivalent before passing to `armcc`. These are shown in Table 7-6.

Table 7-5 Command-line options

--16	--arm_only	--diag_error	--diag_warning	--li
--32	--bi	--diag_remark	--fpu	--library_type
--apcs	--cpu	--diag_style	--fpumode	--thumb
--arm	--device	--diag_suppress	--i	--unaligned_access --no_unaligned_access

Table 7-6 armcc equivalent command-line options

armasm	armcc
--16	--thumb
--32	--arm
--i	--I

To pass other simple compiler options, such as the preprocessor option `-D`, you must use the `--cpreproc_opts` command-line option. `armasm` correctly interprets the preprocessed `#line` commands. It can generate error messages and `debug_line` tables using the information in the `#line` commands.

Example 7-1 on page 7-19 shows the commands you write to preprocess and assemble a file, `source.s`. The example also passes the compiler options to define a macro called `RELEASE`, and to undefine a macro called `ALPHA`.

Example 7-1 Preprocessing an assembly language source file

```
armasm --cpreproc --cpreproc_opts=-D,RELEASE,-U,ALPHA source.s
```

If you want to use complex preprocessor options, you must manually call `armcc` to preprocess the file before calling `armasm`. Example 7-2 shows the commands you write to manually preprocess and assemble a file, `source.s`. In this example, the preprocessor outputs a file called `preprocessed.s`, and `armasm` assembles `preprocessed.s`.

Example 7-2 Preprocessing an assembly language source file manually

```
armcc -E source.s > preprocessed.s  
armasm preprocessed.s
```

7.11.1 See also**Concepts**

Using the Compiler:

- *Compiler command-line options listed by group* on page 3-4.

Reference

Assembler Reference:

- `--cpreproc` on page 2-9
- `--cpreproc_opts=options` on page 2-9.

7.12 Address alignment

For processors based on ARMv5 or earlier, or ARMv6-M, you must ensure that addresses for 4-byte transfers are 4-byte word-aligned, and addresses for 2-byte transfers are 2-byte aligned. In ARMv6 and later, except ARMv6-M, unaligned accesses are permitted for LDR, LDRH, STR, STRH, LDRSH, LDRT, STRT, LDRSHT, LDRHT, STRHT, and TBH instructions, where the architecture supports the instruction.

On some ARM processors, you can enable alignment checking. Non word-aligned 32-bit transfers cause an alignment exception if alignment checking is enabled.

If all your accesses are aligned, you can use the `--no_unaligned_access` command line option, to avoid linking in any library functions that might have an unaligned option.

If a processor does not have alignment checking available and enabled:

- For STR, the specified address is rounded down to a multiple of four.
- For LDR:
 1. The specified address is rounded down to a multiple of four.
 2. Four bytes of data are loaded from the resulting address.
 3. The loaded data is rotated right by one, two or three bytes according to bits [1:0] of the address.

For a little-endian memory system, this causes the addressed byte to occupy the least significant byte of the register.

For a big-endian memory system, it causes the addressed byte to occupy:

- bits[31:24] if bit[0] of the address is 0
- bits[15:8] if bit[0] of the address is 1.

- For STM, LDM, STRD, and LDRD, in ARMv6 and earlier architectures, the specified address is rounded down to a multiple of 4.

In ARMv7 some instructions will fault regardless of alignment checking.

7.12.1 See also

Reference

Assembler Reference:

- `--no_unaligned_access` on page 2-24.

7.13 Instruction width selection in Thumb

If you are writing Thumb code for ARMv6T2 or later processors, some instructions can have either a 16-bit encoding or a 32-bit encoding.

If you do not specify the instruction size, by default:

- for forward reference LDR, ADR, and B instructions, the assembler always generates a 16-bit instruction, even if that results in failure for a target that could be reached using a 32-bit instruction.
- for external reference LDR and B instructions, the assembler always generates a 32-bit instruction.
- in all other cases, the assembler generates the smallest size encoding that can be output.

If you want to over-ride this behavior, you can use the `.W` or `.N` width specifier to ensure a particular instruction size. The assembler will fault if it cannot generate an instruction with the specified width.

The `.W` specifier is ignored when assembling to ARM code, so you can safely use this specifier in code that might assemble to either ARM or Thumb code. However, the `.N` specifier will be faulted when assembling to ARM code.

7.13.1 See also

Reference

Assembler Reference:

- *Instruction width specifiers* on page 3-11.

Chapter 8

Symbols, Literals, Expressions, and Operators

The following topics describe how you can use symbols to represent variables, addresses and constants in code. It also describes how you can combine these with operators to create numeric or string expressions:

- *Symbol naming rules* on page 8-3
- *Variables* on page 8-4
- *Numeric constants* on page 8-6
- *Assembly time substitution of variables* on page 8-7
- *Register-relative and PC-relative expressions* on page 8-9
- *Labels* on page 8-11
- *Labels for PC-relative addresses* on page 8-12
- *Labels for register-relative addresses* on page 8-13
- *Labels for absolute addresses* on page 8-14
- *Local labels* on page 8-15
- *Syntax of local labels* on page 8-17
- *String expressions* on page 8-18
- *String literals* on page 8-19
- *Numeric expressions* on page 8-20
- *Numeric literals* on page 8-21

- *Floating-point literals* on page 8-23
- *Logical expressions* on page 8-25
- *Logical literals* on page 8-26
- *Unary operators* on page 8-27
- *Binary operators* on page 8-29
- *Multiplicative operators* on page 8-30
- *String manipulation operators* on page 8-31
- *Shift operators* on page 8-32
- *Addition, subtraction, and logical operators* on page 8-33
- *Relational operators* on page 8-34
- *Boolean operators* on page 8-35
- *Operator precedence* on page 8-36
- *Difference between operator precedence in armasm and C* on page 8-37.

8.1 Symbol naming rules

The following general rules apply to symbol names:

- Symbol names must be unique within their scope.
- You can use uppercase letters, lowercase letters, numeric characters, or the underscore character in symbol names. Symbol names are case-sensitive, and all characters in the symbol name are significant.
- Do not use numeric characters for the first character of symbol names, except in local labels.
- Symbols must not use the same name as built-in variable names or predefined symbol names.
- If you use the same name as an instruction mnemonic or directive, use double bars to delimit the symbol name. For example:
`||ASSERT||`
 The bars are not part of the symbol.
- You must not use the symbols `$_a`, `$_t`, `$_t.x`, or `$_d` as program labels. These are mapping symbols used to mark ARM, Thumb, ThumbEE, and data within the object file.

If you have to use a wider range of characters in symbols, for example, when working with compilers, use single bars to delimit the symbol name. For example:

`|.text|`

The bars are not part of the symbol. You cannot use bars, semicolons, or newlines within the bars.

8.1.1 See also

Concepts

- *Local labels* on page 8-15
- *Predeclared core register names* on page 3-15
- *Predeclared extension register names* on page 3-16
- *Predeclared XScale register names* on page 3-17
- *Predeclared coprocessor names* on page 3-19
- *Built-in variables and constants* on page 7-7.

8.2 Variables

The value of a variable can be changed as assembly proceeds. Variables are local to the assembler. This means that in the generated code or data, every instance of the variable has a fixed value.

Variables are of three types:

- numeric
- logical
- string.

The type of a variable cannot be changed.

The range of possible values of a numeric variable is the same as the range of possible values of a numeric constant or numeric expression.

The possible values of a logical variable are {TRUE} or {FALSE}.

The range of possible values of a string variable is the same as the range of values of a string expression.

Use the *GBLA*, *GBLL*, *GBLS*, *LCLA*, *LCLL*, and *LCLS* directives to declare symbols representing variables, and assign values to them using the *SETA*, *SETL*, and *SETS* directives.

8.2.1 Example

```

a   SETA 100;
L1  MOV R1, #(a*5); In the object file, this is MOV R1, #500
a   SETA 200;  Value of 'a' is 200 only after this point.
      ;  The previous instruction will always be MOV R1, #500
...
BNE L1;          When the processor branches to L1, it executes MOV R1, #500

```

8.2.2 See also

Concepts

- *Numeric constants* on page 8-6
- *Numeric expressions* on page 8-20
- *String expressions* on page 8-18
- *Logical expressions* on page 8-25.

Reference

Assembler Reference:

- *GBLA*, *GBLL*, and *GBLS* on page 6-5
- *LCLA*, *LCLL*, and *LCLS* on page 6-7

- *SETA, SETL, and SETS* on page 6-8.

8.3 Numeric constants

Numeric constants are 32-bit integers. You can set them using unsigned numbers in the range 0 to $2^{32}-1$, or signed numbers in the range -2^{31} to $2^{31}-1$. However, the assembler makes no distinction between $-n$ and $2^{32}-n$. Relational operators such as $>=$ use the unsigned interpretation. This means that $0 > -1$ is {FALSE}.

Use the EQU directive to define constants. You cannot change the value of a numeric constant after you define it. You can construct expressions by combining numeric constants and binary operators.

8.3.1 See also

Concept

- *Numeric expressions* on page 8-20
- *Numeric literals* on page 8-21.

Reference

Assembler Reference:

- *EQU* on page 6-89.

8.4 Assembly time substitution of variables

You can use a string variable for a whole line of assembly language, or any part of a line. Use the variable with a \$ prefix in the places where the value is to be substituted for the variable. The dollar character instructs the assembler to substitute the string into the source code line before checking the syntax of the line. The assembler faults if the substituted line is larger than the source line limit.

Numeric and logical variables can also be substituted. The current value of the variable is converted to a hexadecimal string (or T or F for logical variables) before substitution.

Use a dot to mark the end of the variable name if the following character would be permissible in a symbol name. You must set the contents of the variable before you can use it.

If you require a \$ that you do not want to be substituted, use \$\$\$. This is converted to a single \$.

You can include a variable with a \$ prefix in a string. Substitution occurs in the same way as anywhere else.

Substitution does not occur within vertical bars, except that vertical bars within double quotes do not affect substitution.

8.4.1 Example

```

; straightforward substitution
GBLS    add4ff
;
add4ff SETS    "ADD    r4,r4,#0xFF"    ; set up add4ff
        $add4ff.00                    ; invoke add4ff
; this produces
        ADD    r4,r4,#0xFF00
; elaborate substitution
GBLS    s1
GBLS    s2
GBLS    fixup
GBLA    count
;
count   SETA    14
s1      SETS    "a$$b$count" ; s1 now has value a$b0000000E
s2      SETS    "abc"
fixup   SETS    "|xy$s2.z|" ; fixup now has value |xyabcz|
|C$$code| MOV    r4,#16      ; but the label here is C$$code

```

8.4.2 See also

Concept

- *Syntax of source lines in assembly language* on page 4-2
- *Symbol naming rules* on page 8-3.

8.5 Register-relative and PC-relative expressions

Addresses can be represented as a register-relative or PC-relative expression.

A register-relative expression evaluates to a named register combined with a numeric expression.

A PC-relative expression is written in source code as the PC or a label combined with a numeric expression. It can also be expressed in the form `[PC, #number]`. It is represented in the instruction as the PC value plus or minus a numeric offset. The assembler calculates the required offset from the label and the address of the current instruction. If the offset is too big, the assembler produces an error.

It is recommended to write PC-relative expressions using labels rather than PC because the value of PC depends on the instruction set.

Note

- In ARM state, the value of the PC is the address of the current instruction plus 8 bytes.
 - In Thumb state:
 - For B, BL, CBNZ, and CBZ instructions, the value of the PC is the address of the current instruction plus 4 bytes.
 - For all other instructions that use labels, the value of the PC is the address of the current instruction plus 4 bytes, with bit[1] of the result cleared to 0 to make it word-aligned.
-

8.5.1 Example

```

                                LDR    r4,=data+4*n    ; n is an assembly-time variable
                                ; code
                                MOV    pc,lr
data DCD    value_0
                                ; n-1 DCD directives
                                DCD    value_n          ; data+4*n points here
                                ; more DCD directives

```

8.5.2 See also

Concepts

- *Labels* on page 8-11.

Reference

Assembler Reference:

- *MAP* on page 6-21.

8.6 Labels

Labels are symbols representing the memory addresses of instructions or data. The address can be PC-relative, register-relative, or absolute. Labels are local to the source file unless you make them global using the `EXPORT` directive.

The address given by a label is calculated during assembly. The assembler calculates the address of a label relative to the origin of the section where the label is defined. A reference to a label within the same section can use the PC plus or minus an offset. This is called *PC-relative addressing*.

Addresses of labels in other sections are calculated at link time, when the linker has allocated specific locations in memory for each section.

8.6.1 See also

Concept

- *Labels for PC-relative addresses* on page 8-12
- *Labels for register-relative addresses* on page 8-13
- *Labels for absolute addresses* on page 8-14.

Reference

Assembler Reference:

- *EXPORT or GLOBAL* on page 6-91.

8.7 Labels for PC-relative addresses

These represent the PC, plus or minus a numeric value. Use them as targets for branch instructions, or to access small items of data embedded in code sections. You can define PC-relative labels using a label on an instruction or on one of the data definition directives.

You can also use the section name of an AREA directive as a label for PC-relative addresses. In this case the label points to the first byte of the specified AREA. Using AREA names as branch targets is not recommended because when branching from ARM to Thumb state or Thumb to ARM state in this way, the processor does not change the state properly.

8.7.1 See also

Reference

Assembler Reference:

- *AREA* on page 6-81
- *DCB* on page 6-25
- *DCD and DCDU* on page 6-26
- *DCFD and DCFDU* on page 6-29
- *DCFS and DCFSU* on page 6-31
- *DCI* on page 6-32
- *DCQ and DCQU* on page 6-34
- *DCW and DCWU* on page 6-36.

8.8 Labels for register-relative addresses

These represent a named register plus a numeric value. They are most often used to access data in data sections. You can define them with a storage map. You can use the EQU directive to define additional register-relative labels, based on labels defined in storage maps.

Example 8-1 Storage map definitions

MAP	0, r9
MAP	0xff, r9

8.8.1 See also

Reference

Assembler Reference:

- *MAP* on page 6-21
- *SPACE or FILL* on page 6-23
- *DCDO* on page 6-28
- *EQU* on page 6-89.

8.9 Labels for absolute addresses

These are numeric constants. They are integers in the range 0 to $2^{32}-1$. They address the memory directly. You can use labels to represent absolute addresses using the EQU directive. You can specify the absolute address as ARM, Thumb, or data to ensure that the labels are used correctly when referenced in code.

Example 8-2 Defining labels for absolute address

```

abc EQU 2           ; assigns the value 2 to the symbol abc.
xyz EQU label+8      ; assigns the address (label+8) to the
                    ; symbol xyz.
fiq EQU 0x1C, CODE32 ; assigns the absolute address 0x1C to
                    ; the symbol fiq, and marks it as code

```

8.9.1 See also

Concepts

- *Labels* on page 8-11
- *Labels for PC-relative addresses* on page 8-12
- *Labels for register-relative addresses* on page 8-13.

Reference

Assembler Reference:

- *EQU* on page 6-89.

8.10 Local labels

Local labels are a subclass of label. A local label is a number in the range 0-99, optionally followed by a name. Unlike other labels, a local label can be defined many times and the same number can be used for more than one local label in an area.

Local labels do not appear in the object file. This means that, for example, a debugger cannot set a breakpoint directly on a local label, like it can for labels kept using the KEEP directive.

A local label can be used in place of *symbol* in source lines in an assembly language module:

- on its own, that is, where there is no instruction or directive
- on a line that contains an instruction
- on a line that contains a code- or data-generating directive.

A local label is generally used where you might use a PC-relative label.

Local labels are typically used for loops and conditional code within a routine, or for small subroutines that are only used locally. They are particularly useful when you are generating labels in macros.

The scope of local labels is limited by the AREA directive. Use the ROUT directive to limit the scope of local labels more tightly. A reference to a local label refers to a matching label within the same scope. If there is no matching label within the scope in either direction, the assembler generates an error message and the assembly fails.

You can use the same number for more than one local label even within the same scope. By default, the assembler links a local label reference to:

- the most recent local label of the same number, if there is one within the scope
- the next following local label of the same number, if there is not a preceding one within the scope.

Use the optional parameters to modify this search pattern if required.

8.10.1 See also

Concepts

- *Syntax of local labels* on page 8-17
- *Labels* on page 8-11
- *Syntax of source lines in assembly language* on page 4-2.

Reference

Assembler Reference:

- *MACRO and MEND* on page 6-40
- *KEEP* on page 6-99
- *ROUT* on page 6-102.

8.11 Syntax of local labels

The syntax of a local label is:

n{*rouname*}

The syntax of a reference to a local label is:

%{F|B}{A|T}*n*{*rouname*}

where:

<i>n</i>	is the number of the local label in range 0-99.
<i>rouname</i>	is the name of the current scope.
%	introduces the reference.
F	instructs the assembler to search forwards only.
B	instructs the assembler to search backwards only.
A	instructs the assembler to search all macro levels.
T	instructs the assembler to look at this macro level only.

If neither F nor B is specified, the assembler searches backwards first, then forwards.

If neither A nor T is specified, the assembler searches all macros from the current level to the top level, but does not search lower level macros.

If *rouname* is specified in either a label or a reference to a label, the assembler checks it against the name of the nearest preceding *ROUT* directive. If it does not match, the assembler generates an error message and the assembly fails.

8.11.1 See also

Concepts

- *Local labels* on page 8-15.

Reference

Assembler Reference:

- *ROUT* on page 6-102.

8.12 String expressions

String expressions consist of combinations of string literals, string variables, string manipulation operators, and parentheses.

Characters that cannot be placed in string literals can be placed in string expressions using the `:CHR:` unary operator. Any ASCII character from 0 to 255 is permitted.

The value of a string expression cannot exceed 5120 characters in length. It can be of zero length.

8.12.1 Example

```
improb SETS    "literal":CC:(strvar2:LEFT:4)
               ; sets the variable improb to the value "literal"
               ; with the left-most four characters of the
               ; contents of string variable strvar2 appended
```

8.12.2 See also

Concepts

- *Variables* on page 8-4
- *String literals* on page 8-19
- *Unary operators* on page 8-27
- *String manipulation operators* on page 8-31.

Reference

Assembler Reference:

- *SETA, SETL, and SETS* on page 6-8.

8.13 String literals

String literals consist of a series of characters or spaces contained between double quote characters. The length of a string literal is restricted by the length of the input line.

To include a double quote character or a dollar character within the string literal, include the character twice as a pair. For example, you must use \$\$ if you require a single \$ in the string.

C string escape sequences are also enabled and can be used within the string, unless `--no_esc` is specified.

8.13.1 Examples

```
abc    SETS    "this string contains only one "" double quote"
def    SETS    "this string contains only one $$ dollar symbol"
```

8.13.2 See also

Concepts

- *Syntax of source lines in assembly language* on page 4-2.

Reference

Assembler Reference:

- `--no_esc` on page 2-22.

8.14 Numeric expressions

Numeric expressions consist of combinations of numeric constants, numeric variables, ordinary numeric literals, binary operators, and parentheses.

Numeric expressions can contain register-relative or program-relative expressions if the overall expression evaluates to a value that does not include a register or the PC.

Numeric expressions evaluate to 32-bit integers. You can interpret them as unsigned numbers in the range 0 to $2^{32}-1$, or signed numbers in the range -2^{31} to $2^{31}-1$. However, the assembler makes no distinction between $-n$ and $2^{32}-n$. Relational operators such as $>=$ use the unsigned interpretation. This means that $0 > -1$ is {FALSE}.

8.14.1 Example

```
a  SETA    256*256          ; 256*256 is a numeric expression
   MOV     r1,#(a*22)      ; (a*22) is a numeric expression
```

8.14.2 See also

Concepts

- *Numeric constants* on page 8-6
- *Variables* on page 8-4
- *Numeric literals* on page 8-21
- *Binary operators* on page 8-29.

Reference

Assembler Reference:

- *SETA, SETL, and SETS* on page 6-8.

8.15 Numeric literals

Numeric literals can take any of the following forms:

decimal-digits

0xhexadecimal-digits

&hexadecimal-digits

n_base-n-digits

'character'

where:

decimal-digits Is a sequence of characters using only the digits 0 to 9.

hexadecimal-digits Is a sequence of characters using only the digits 0 to 9 and the letters A to F or a to f.

n_ Is a single digit between 2 and 9 inclusive, followed by an underscore character.

base-n-digits Is a sequence of characters using only the digits 0 to ($n-1$)

character Is any single character except a single quote. Use the standard C escape character (\) if you require a single quote. The character must be enclosed within opening and closing single quotes. In this case the value of the numeric literal is the numeric code of the character.

You must not use any other characters. The sequence of characters must evaluate to an integer in the range 0 to $2^{32}-1$ (except in DCQ and DCQU directives, where the range is 0 to $2^{64}-1$).

8.15.1 Examples

a	SETA	34906	
addr	DCD	0xA10E	
	LDR	r4,=&1000000F	
	DCD	2_11001010	
c3	SETA	8_74007	
	DCQ	0x0123456789abcdef	
	LDR	r1,='A'	; pseudo-instruction loading 65 into r1
	ADD	r3,r2,#'\''	; add 39 to contents of r2, result to r3

8.15.2 See also

Concepts

- *Numeric constants* on page 8-6.

8.16 Floating-point literals

Floating-point literals can take any of the following forms:

```
{-}digitsE{-}digits
{-}{digits}.digits
{-}{digits}.digitsE{-}digits
0xhexdigits
&hexdigits
0f_hexdigits
0d_hexdigits
```

where:

- | | |
|------------------|---|
| <i>digits</i> | Are sequences of characters using only the digits 0 to 9. You can write E in uppercase or lowercase. These forms correspond to normal floating-point notation. |
| <i>hexdigits</i> | Are sequences of characters using only the digits 0 to 9 and the letters A to F or a to f. These forms correspond to the internal representation of the numbers in the computer. Use these forms to enter infinities and NaNs, or if you want to be sure of the exact bit patterns you are using. |

The 0x and & forms allow the floating-point bit pattern to be specified by any number of hex digits.

The 0f_ form requires the floating-point bit pattern to be specified by exactly 8 hex digits.

The 0d_ form requires the floating-point bit pattern to be specified by exactly 16 hex digits.

The range for single-precision floating-point values is:

- maximum 3.40282347e+38
- minimum 1.17549435e-38.

The range for double-precision floating-point values is:

- maximum 1.79769313486231571e+308
- minimum 2.22507385850720138e-308.

Floating-point numbers are only available if your system has VFP, or NEON with floating-point.

8.16.1 Examples

DCFD	1E308,-4E-100	
DCFS	1.0	
DCFS	0.02	
DCFD	3.725e15	
DCFS	0x7FC00000	; Quiet NaN
DCFD	&FFF0000000000000	; Minus infinity

8.16.2 See also

Concepts

- *Numeric constants* on page 8-6
- *Numeric literals* on page 8-21.

8.17 Logical expressions

Logical expressions consist of combinations of logical literals ({TRUE} or {FALSE}), logical variables, Boolean operators, relations, and parentheses.

Relations consist of combinations of variables, literals, constants, or expressions with appropriate relational operators.

8.17.1 See also

Concepts

- *Boolean operators* on page 8-35
- *Relational operators* on page 8-34.

8.18 Logical literals

The logical or boolean literals can have one of two values:

- {TRUE}
- {FALSE}.

8.18.1 See also

Concepts

- *Numeric literals* on page 8-21
- *String literals* on page 8-19.

8.19 Unary operators

Unary operators have the highest precedence and are evaluated first. A unary operator precedes its operand. Adjacent operators are evaluated from right to left.

Table 8-1 lists the unary operators that return strings.

Table 8-1 Unary operators that return strings

Operator	Usage	Description
:CHR:	:CHR:A	Returns the character with ASCII code A.
:LOWERCASE:	:LOWERCASE:string	Returns the given string, with all uppercase characters converted to lowercase.
:REVERSE_CC:	:REVERSE_CC:cond_code	Returns the inverse of the condition code in cond_code, or an error if cond_code does not contain a valid condition code.
:STR:	:STR:A	Returns an 8-digit hexadecimal string corresponding to a numeric expression, or the string "T" or "F" if used on a logical expression.
:UPPERCASE:	:UPPERCASE:string	Returns the given string, with all lowercase characters converted to uppercase.

Table 8-2 lists the unary operators that return numeric values.

Table 8-2 Unary operators that return numeric or logical values

Operator	Usage	Description
?	?A	Number of bytes of executable code generated by line defining symbol A.
+ and -	+A -A	Unary plus. Unary minus. + and – can act on numeric and PC-relative expressions.
:BASE:	:BASE:A	If A is a PC-relative or register-relative expression, :BASE: returns the number of its register component. :BASE: is most useful in macros.
:CC_ENCODING:	:CC_ENCODING:cond_code	Returns the numeric value of the condition code in cond_code, or an error if cond_code does not contain a valid condition code.
:DEF:	:DEF:A	{TRUE} if A is defined, otherwise {FALSE}.
:INDEX:	:INDEX:A	If A is a register-relative expression, :INDEX: returns the offset from that base register. :INDEX: is most useful in macros.
:LEN:	:LEN:A	Length of string A.

Table 8-2 Unary operators that return numeric or logical values (continued)

Operator	Usage	Description
:LNOT:	:LNOT:A	Logical complement of A.
:NOT:	:NOT:A	Bitwise complement of A (~ is an alias, for example ~A).
:RCONST:	:RCONST:Rn	Number of register, 0-15 corresponding to R0-R15.

8.19.1 See also

Concepts

- *Binary operators* on page 8-29.

8.20 Binary operators

Binary operators are written between the pair of sub-expressions they operate on.

Binary operators have lower precedence than unary operators. Binary operators appear in this section in order of precedence.

Note

The order of precedence is not the same as in C.

8.20.1 See also

Concepts

- *Multiplicative operators* on page 8-30
- *String manipulation operators* on page 8-31
- *Shift operators* on page 8-32
- *Addition, subtraction, and logical operators* on page 8-33
- *Relational operators* on page 8-34
- *Boolean operators* on page 8-35
- *Difference between operator precedence in armasm and C* on page 8-37.

8.21 Multiplicative operators

Multiplicative operators have the highest precedence of all binary operators. They act only on numeric expressions.

Table 8-3 shows the multiplicative operators.

Table 8-3 Multiplicative operators

Operator	Alias	Usage	Explanation
*		A*B	Multiply
/		A/B	Divide
:MOD:	%	A:MOD:B	A modulo B

You can use the :MOD: operator on PC-relative expressions in the form of *PC-relative:MOD:Constant*. This enables easier code alignment checks in assembler. For example:

```
AREA x, CODE
ASSERT ({PC}:MOD:4) == 0
DCB 1
y DCB 2
ASSERT (y:MOD:4) == 1
ASSERT ({PC}:MOD:4) == 2
END
```

8.21.1 See also

Concepts

- *Register-relative and PC-relative expressions* on page 8-9
- *Binary operators* on page 8-29
- *Numeric literals* on page 8-21
- *Numeric expressions* on page 8-20.

8.22 String manipulation operators

Table 8-4 shows the string manipulation operators. In CC, both A and B must be strings. In the slicing operators LEFT and RIGHT:

- A must be a string
- B must be a numeric expression.

Table 8-4 String manipulation operators

Operator	Usage	Explanation
:CC:	A:CC:B	B concatenated onto the end of A
:LEFT:	A:LEFT:B	The left-most B characters of A
:RIGHT:	A:RIGHT:B	The right-most B characters of A

8.22.1 See also

Concepts

- *String expressions* on page 8-18
- *Numeric expressions* on page 8-20.

8.23 Shift operators

Shift operators act on numeric expressions, shifting or rotating the first operand by the amount specified by the second.

Table 8-5 shows the shift operators.

Table 8-5 Shift operators

Operator	Alias	Usage	Explanation
:ROL:		A:ROL:B	Rotate A left by B bits
:ROR:		A:ROR:B	Rotate A right by B bits
:SHL:	<<	A:SHL:B	Shift A left by B bits
:SHR:	>>	A:SHR:B	Shift A right by B bits

Note

SHR is a logical shift and does not propagate the sign bit.

8.23.1 See also

Concepts

- *Binary operators* on page 8-29.

8.24 Addition, subtraction, and logical operators

Addition and subtraction operators act on numeric expressions.

Logical operators act on numeric expressions. The operation is performed *bitwise*, that is, independently on each bit of the operands to produce the result.

Table 8-6 shows addition, subtraction, and logical operators.

Table 8-6 Addition, subtraction, and logical operators

Operator	Alias	Usage	Explanation
+		A+B	Add A to B
-		A-B	Subtract B from A
:AND:	&	A:AND:B	Bitwise AND of A and B
:EOR:	^	A:EOR:B	Bitwise Exclusive OR of A and B
:OR:		A:OR:B	Bitwise OR of A and B

The use of | is deprecated. You must use :OR: instead.

8.24.1 See also

Concepts

- *Binary operators* on page 8-29.

8.25 Relational operators

Table 8-7 shows the relational operators. These act on two operands of the same type to produce a logical value.

The operands can be one of:

- numeric
- PC-relative
- register-relative
- strings.

Strings are sorted using ASCII ordering. String A is less than string B if it is a leading substring of string B, or if the left-most character in which the two strings differ is less in string A than in string B.

Arithmetic values are unsigned, so the value of 0>-1 is {FALSE}.

Table 8-7 Relational operators

Operator	Alias	Usage	Explanation
=	==	A=B	A equal to B
>		A>B	A greater than B
>=		A>=B	A greater than or equal to B
<		A<B	A less than B
<=		A<=B	A less than or equal to B
/=	<> !=	A/=B	A not equal to B

8.25.1 See also

Concepts

- *Binary operators* on page 8-29.

8.26 Boolean operators

These are the operators with the lowest precedence. They perform the standard logical operations on their operands.

In all three cases both A and B must be expressions that evaluate to either {TRUE} or {FALSE}.

Table 8-8 shows the Boolean operators.

Table 8-8 Boolean operators

Operator	Alias	Usage	Explanation
:LAND:	&&	A:LAND:B	Logical AND of A and B
:LEOR:		A:LEOR:B	Logical Exclusive OR of A and B
:LOR:		A:LOR:B	Logical OR of A and B

8.26.1 See also

Concepts

- *Binary operators* on page 8-29.

8.27 Operator precedence

The assembler includes an extensive set of operators for use in expressions. Many of the operators resemble their counterparts in high-level languages such as C.

There is a strict order of precedence in their evaluation:

1. Expressions in parentheses are evaluated first.
2. Operators are applied in precedence order.
3. Adjacent unary operators are evaluated from right to left.
4. Binary operators of equal precedence are evaluated from left to right.

8.27.1 See also

Concepts

- *Unary operators* on page 8-27
- *Binary operators* on page 8-29
- *Multiplicative operators* on page 8-30
- *String manipulation operators* on page 8-31
- *Shift operators* on page 8-32
- *Addition, subtraction, and logical operators* on page 8-33
- *Relational operators* on page 8-34
- *Boolean operators* on page 8-35
- *Difference between operator precedence in armasm and C* on page 8-37.

8.28 Difference between operator precedence in armasm and C

The assembler order of precedence is not exactly the same as in C.

For example, $(1 + 2 :SHR: 3)$ evaluates as $(1 + (2 :SHR: 3)) = 1$ in armasm. The equivalent expression in C evaluates as $((1 + 2) >> 3) = 0$.

You are recommended to use brackets to make the precedence explicit.

If your code contains an expression that would parse differently in C, and you are not using the `--unsafe` option, armasm normally gives a warning:

A1466W: Operator precedence means that expression would evaluate differently in C

Table 8-9 shows the order of precedence of operators in armasm, and a comparison with the order in C (see Table 8-10 on page 8-38).

From these tables:

- The highest precedence operators are at the top of the list.
- The highest precedence operators are evaluated first.
- Operators of equal precedence are evaluated from left to right.

Table 8-9 Operator precedence in armasm

armasm precedence	equivalent C operators
unary operators	unary operators
* / :MOD:	* / %
string manipulation	n/a
:SHL: :SHR: :ROR: :ROL:	<< >>
+ - :AND: :OR: :EOR:	+ - & ^
= > >= < <= /= <>	== > >= < <= !=
:LAND: :LOR: :LEOR:	&&

Table 8-10 Operator precedence in C

C precedence
unary operators
* / %
+ - (as binary operators)
<< >>
< <= > >=
== !=
&
^
&&

8.28.1 See also

Concepts

- *Operator precedence* on page 8-36.

Chapter 9

NEON and VFP Programming

This describes the assembly programming of NEON and the VFP coprocessor:

- *Architecture support for NEON and VFP* on page 9-3
- *Half-precision extension* on page 9-4
- *Fused Multiply-Add extension* on page 9-5
- *Extension register bank mapping* on page 9-6
- *NEON views of the register bank* on page 9-8
- *VFP views of the extension register bank* on page 9-9
- *Load values to VFP and NEON registers* on page 9-10
- *Conditional execution of NEON and VFP instructions* on page 9-11
- *Floating-point exceptions* on page 9-12
- *NEON and VFP data types* on page 9-14
- *NEON vectors* on page 9-16
- *Normal NEON instructions* on page 9-17
- *Long NEON instructions* on page 9-18
- *Wide NEON instructions* on page 9-19
- *Narrow NEON instructions* on page 9-20
- *Saturating NEON instructions* on page 9-21
- *NEON scalars* on page 9-22

- *Extended notation* on page 9-23
- *Polynomial arithmetic over $\{0,1\}$* on page 9-24
- *NEON and VFP system registers* on page 9-25
- *FPSCR, the floating-point status and control register* on page 9-26
- *FPExc, the floating-point exception register* on page 9-29
- *FPSID, the floating-point system ID register* on page 9-30
- *Flush-to-zero mode* on page 9-31
- *When to use flush-to-zero mode* on page 9-32
- *The effects of using flush-to-zero mode* on page 9-33
- *Operations not affected by flush-to-zero mode* on page 9-34
- *VFP vector mode* on page 9-35
- *Vectors in the VFP extension register bank* on page 9-36
- *VFP vector wrap-around* on page 9-38
- *VFP vector stride* on page 9-39
- *Restriction on vector length* on page 9-40
- *Control of scalar, vector, and mixed operations* on page 9-41
- *VFP directives and vector notation* on page 9-43
- *Pre-UAL VFP mnemonics* on page 9-44
- *Vector notation* on page 9-46
- *VFPASSERT SCALAR* on page 9-47
- *VFPASSERT VECTOR* on page 9-48.

9.1 Architecture support for NEON and VFP

The NEON extension is optionally available only for the ARMv7-A and ARMv7-R architectures, and is available for the ARMv7E-M architecture. All NEON instructions, with the exception of half-precision and fused multiply-add instructions, are available on systems that support NEON. Some of these instructions are also available on systems that implement VFP extension without NEON. These are called shared instructions.

Most VFP and the shared instructions are available in all versions of the VFP architecture. Where this is not true, the descriptions of the instructions specify the applicable VFP architecture versions.

VFPv3 has variants that do not support all VFPv3 registers and floating-point data types. VFPv3 with half-precision extension and fused multiply-add extension is called VFPv4. For details of the implemented VFP architecture and variant, you must always refer to the appropriate product documentation. To get VFP, you must specify the FPU or have it implicit in the CPU.

NEON and VFP instructions, including the half-precision and fused multiply-add instructions, are treated as Undefined Instructions on systems that do not support the necessary architecture extension. Even on systems that support NEON and VFP, the instructions are undefined if the necessary coprocessors are not enabled in the Coprocessor Access Control Register (CP15 CPACR).

9.1.1 See also

Concepts

- *Half-precision extension* on page 9-4
- *Fused Multiply-Add extension* on page 9-5
- *VFP vector mode* on page 9-35.

Using ARM Libraries:

- Chapter 4 *Floating-point support*.

Reference

Technical Reference Manual for your processor.

9.2 Half-precision extension

The Half-precision extension is an optional architecture that extends both the VFPv3 and the NEON architectures. It provides VFP and NEON instructions that perform conversion between single-precision (32-bit) and half-precision (16-bit) floating-point numbers.

The half-precision instructions are only available on NEON or VFP systems that implement the half-precision extension. The VFP variants that implement the half-precision extension are VFPv3-FP16, VFPv3-D16-FP16, and VFPv4.

9.2.1 See also

Concepts

- *Architecture support for NEON and VFP* on page 9-3.

9.3 Fused Multiply-Add extension

The Fused Multiply-Add extension is an optional architecture that extends both the VFPv3 and the NEON architectures. It provides VFP and NEON instructions that perform multiply and accumulate operations with a single rounding step, so suffers from less loss of accuracy than performing a multiplication followed by an add.

The fused multiply-add instructions are only available on NEON or VFP systems that implement the fused multiply-add extension. The VFP system that implements the fused multiply-add extension is VFPv4.

9.3.1 See also

Concepts

- *Architecture support for NEON and VFP* on page 9-3.

9.4 Extension register bank mapping

NEON and VFP use the same extension register bank. This is distinct from the ARM register bank. The extension register bank is a collection of registers which can be accessed as either 32-bit, 64-bit, or 128-bit registers, depending on whether the instruction is NEON or VFP.

Figure 9-1 on page 9-7 shows the three views of the extension register bank, and the overlap between the different size registers. For example, the 128-bit register Q0 is an alias for two consecutive 64-bit registers D0 and D1, and is also an alias for four consecutive 32-bit registers S0, S1, S2, and S3. The 128-bit register Q8 is an alias for 2 consecutive 64-bit registers D16 and D17 but does not have an alias using the 32-bit "Sn" registers.

Note

If your processor has both NEON and VFP, all the NEON registers overlap with the VFP registers.

The aliased views enables half-precision, single-precision, and double-precision values, and NEON vectors to coexist in different non-overlapped registers at the same time.

You can also use the same overlapped registers to store half-precision, single-precision, and double-precision values, and NEON vectors at different times.

Do not attempt to use overlapped 32-bit and 64-bit, or 128-bit registers at the same time because it creates meaningless results.

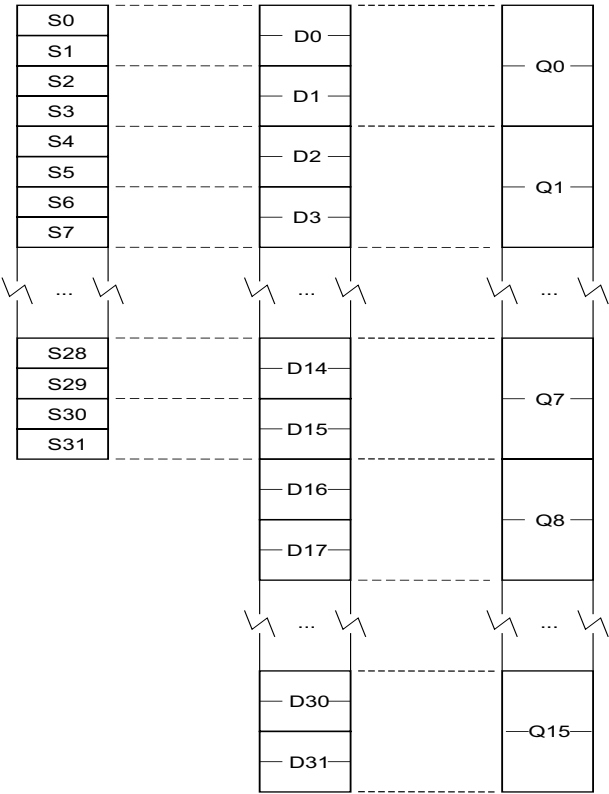


Figure 9-1 Extension register bank

The mapping between the registers is as follows:

- S<2n> maps to the least significant half of D<n>
- S<2n+1> maps to the most significant half of D<n>
- D<2n> maps to the least significant half of Q<n>
- D<2n+1> maps to the most significant half of Q<n>.

For example, you can access the least significant half of the elements of a vector in Q6 by referring to D12, and the most significant half of the elements by referring to D13.

9.4.1 See also

Concepts

- NEON views of the register bank on page 9-8
- VFP views of the extension register bank on page 9-9.

9.5 NEON views of the register bank

In NEON, you can view the extension register bank as:

- Sixteen 128-bit registers, Q0-Q15.
- Thirty-two 64-bit registers, D0-D31.
- A combination of registers from the above views.

NEON views each register as containing a *vector* of 1, 2, 4, 8, or 16 elements, all of the same size and type. Individual elements can also be accessed as *scalars*.

In NEON, the 64-bit registers are called doubleword registers and the 128-bit registers are called quadword registers.

9.5.1 See also

Concepts

- *VFP views of the extension register bank* on page 9-9
- *Extension register bank mapping* on page 9-6.

9.6 VFP views of the extension register bank

In VFPv3 and VFPv3-FP16, you can view the extension register bank as:

- Thirty-two 64-bit registers, D0-D31.
- Thirty-two 32-bit registers, S0-S31. Only half of the register bank is accessible in this view.
- A combination of registers from the above views.

In VFPv2, VFPv3-D16, and VFPv3-D16-FP16, you can view the extension register bank as:

- Sixteen 64-bit registers, D0-D15.
- Thirty-two 32-bit registers, S0-S31.
- A combination of registers from the above views.

In VFP, 64-bit registers are called double-precision registers and can contain double-precision floating-point values. 32-bit registers are called single-precision registers and can contain either a single-precision or two half-precision floating-point values.

9.6.1 See also

Concepts

- *NEON views of the register bank* on page 9-8
- *Extension register bank mapping* on page 9-6.

9.7 Load values to VFP and NEON registers

In the NEON and VFPv3 instruction sets, there are instructions to load a limited range of floating-point immediate values.

The NEON *VMOV* and *VMVN* instructions can also load integer values.

You can load any 64-bit integer, single-precision, or double-precision floating-point value from a literal pool, in a single instruction, using the *VLDR* pseudo-instruction.

9.7.1 See also

Reference

Assembler Reference:

- *VMOV*, *VMVN (immediate)* on page 4-34
- *VMOV* on page 4-106
- *VLDR pseudo-instruction* on page 4-83.

9.8 Conditional execution of NEON and VFP instructions

In ARM state, you can use a condition code to control the execution of VFP instructions. The instruction is executed conditionally, according to the status flags in the APSR, in exactly the same way as almost all other ARM instructions.

In ARM state, except for the instructions that are common to both VFP and NEON, you cannot use a condition code to control the execution of NEON instructions.

In Thumb state on a Thumb-2 processor, you can use an IT instruction to set condition codes on up to four following NEON or VFP instructions.

9.8.1 See also

Concepts

- *Condition code meanings* on page 6-9.

9.9 Floating-point exceptions

The Advanced SIMD and VFP extensions record the following floating-point exceptions in the FPSCR cumulative flags:

Invalid operation

The exception is caused if the result of an operation has no mathematical value or cannot be represented.

Division by zero

The exception is caused if a divide operation has a zero divisor and a dividend that is not zero, an infinity or a NaN.

Overflow

The exception is caused if the absolute value of the result of an operation, produced after rounding, is greater than the maximum positive normalized number for the destination precision.

Underflow

The exception is caused if the absolute value of the result of an operation, produced before rounding, is less than the minimum positive normalized number for the destination precision, and the rounded result is inexact.

Inexact

The exception is caused if the result of an operation is not equivalent to the value that would be produced if the operation were performed with unbounded precision and exponent range.

Input denormal

The exception is caused if a denormalized input operand is replaced in the computation by a zero.

In the *Assembler Reference*, in the descriptions of the instructions that can cause floating-point exceptions, there is a subsection listing the exceptions. If there is no Floating-point exceptions subsection in an instruction description, that instruction cannot cause any floating-point exception.

9.9.1 See also

Concepts

- *Flush-to-zero mode* on page 9-31
- *FPSCR, the floating-point status and control register* on page 9-26.

Reference

- The *Technical Reference Manual* for your VFP coprocessor.

9.10 NEON and VFP data types

Data type specifiers in NEON and VFP instructions consist of a letter indicating the type of data, usually followed by a number indicating the width. They are separated from the instruction mnemonic by a point. Table 9-1 shows the data types available in NEON instructions. Table 9-2 shows the data types available in VFP instructions.

Table 9-1 NEON data types

	8-bit	16-bit	32-bit	64-bit
Unsigned integer	U8	U16	U32	U64
Signed integer	S8	S16	S32	S64
Integer of unspecified type	I8	I16	I32	I64
Floating-point number	not available	F16	F32 (or F)	not available
Polynomial over {0,1}	P8	P16	not available	not available

Table 9-2 VFP data types

	16-bit	32-bit	64-bit
Unsigned integer	U16	U32	not available
Signed integer	S16	S32	not available
Floating-point number	F16	F32 (or F)	F64 (or D)

The datatype of the second (or only) operand is specified in the instruction.

———— **Note** ————

- Most instructions have a restricted range of permitted data types. See the instruction pages for details. However, the data type description is flexible:
 - If the description specifies I, you can also use S or U data types
 - If only the data size is specified, you can specify a type (I, S, U, P or F)
 - If no data type is specified, you can specify a data type.
- The F16 data type is only available on systems that implement the half-precision architecture extension.

9.10.1 See also

Concepts

- *Polynomial arithmetic over $\{0,1\}$* on page 9-24.

9.11 NEON vectors

A NEON operand can be a vector or a scalar. A NEON vector can be a 64-bit doubleword vector or a 128-bit quadword vector.

The size of the elements in a NEON vector is specified by a datatype suffix in the NEON instruction.

Doubleword vectors can contain:

- eight 8-bit elements
- four 16-bit elements
- two 32-bit elements
- one 64-bit element.

Quadword vectors can contain:

- sixteen 8-bit elements
- eight 16-bit elements
- four 32-bit elements
- two 64-bit elements.

9.11.1 See also

Concepts

- *NEON scalars* on page 9-22
- *Extension register bank mapping* on page 9-6
- *Extended notation* on page 9-23
- *NEON and VFP data types* on page 9-14.

9.12 Normal NEON instructions

Many NEON data processing instructions are available in Normal, Long, Wide, Narrow, and saturating variants.

Normal instructions can operate on any of these vector types, and produce result vectors the same size, and usually the same type, as the operand vectors.

You can specify that the operands and result of a normal instruction must all be quadwords by appending a Q to the instruction mnemonic. If you do this, the assembler produces an error if the operands or result are not quadwords.

9.12.1 See also

Concepts

- *NEON vectors* on page 9-16
- *Long NEON instructions* on page 9-18
- *Wide NEON instructions* on page 9-19
- *Narrow NEON instructions* on page 9-20
- *Saturating NEON instructions* on page 9-21.

9.13 Long NEON instructions

Long instructions operate on doubleword vector operands and produce a quadword vector result. The elements of the result are usually twice the width of the elements of the operands, and the same type.

Long instructions are specified using an L appended to the instruction mnemonic.

9.13.1 See also

Concepts

- *NEON vectors* on page 9-16
- *Normal NEON instructions* on page 9-17
- *Wide NEON instructions* on page 9-19
- *Narrow NEON instructions* on page 9-20
- *Saturating NEON instructions* on page 9-21.

9.14 Wide NEON instructions

Wide instructions operate on one doubleword vector operand and one quadword vector operand. They produce a quadword vector result. The elements of the result and the first operand are twice the width of the elements of the second operand.

Wide instructions are specified using a *W* appended to the instruction mnemonic.

9.14.1 See also

Concepts

- *NEON vectors* on page 9-16
- *Normal NEON instructions* on page 9-17
- *Long NEON instructions* on page 9-18
- *Narrow NEON instructions* on page 9-20
- *Saturating NEON instructions* on page 9-21.

9.15 Narrow NEON instructions

Narrow instructions operate on quadword vector operands, and produce a doubleword vector result. The elements of the result are half the width of the elements of the operands.

Narrow instructions are specified using an N appended to the instruction mnemonic.

9.15.1 See also

Concepts

- *NEON vectors* on page 9-16
- *Normal NEON instructions* on page 9-17
- *Long NEON instructions* on page 9-18
- *Wide NEON instructions* on page 9-19
- *Saturating NEON instructions* on page 9-21.

9.16 Saturating NEON instructions

Saturating instructions saturate the result to the value of the upper limit or lower limit if the result overflows or underflows. The saturation limits depend on the datatype of the instruction. See Table 9-3 for the ranges that NEON saturating instructions saturate to, where x is the result of the operation.

Table 9-3 NEON saturation ranges

Data type	Saturation range of x
S8	$-2^7 \leq x < 2^7$
S16	$-2^{15} \leq x < 2^{15}$
S32	$-2^{31} \leq x < 2^{31}$
S64	$-2^{63} \leq x < 2^{63}$
U8	$0 \leq x < 2^8$
U16	$0 \leq x < 2^{16}$
U32	$0 \leq x < 2^{32}$
U64	$0 \leq x < 2^{64}$

Saturating instructions are specified using a Q prefix between the V and the instruction mnemonic.

9.16.1 See also

Concepts

- *NEON vectors* on page 9-16
- *Normal NEON instructions* on page 9-17
- *Long NEON instructions* on page 9-18
- *Wide NEON instructions* on page 9-19
- *Narrow NEON instructions* on page 9-20.

Reference

Assembler Reference:

- *Saturating instructions* on page 3-107.

9.17 NEON scalars

Some NEON instructions act on scalars in combination with vectors. NEON scalars can be 8-bit, 16-bit, 32-bit, or 64-bit. The instruction syntax refers to the scalars using an index, x , into a doubleword vector, so that $Dm[x]$ is the x th element in vector Dm . Other than multiply instructions, instructions that access scalars can access any element in the register bank.

Multiply instructions only allow 16-bit or 32-bit scalars, and can only access the first 32 scalars in the register bank. That is, in multiply instructions:

- 16-bit scalars are restricted to registers D0-D7, with x in the range 0-3
- 32-bit scalars are restricted to registers D0-D15, with x either 0 or 1.

9.17.1 See also

Concepts

- *Extension register bank mapping* on page 9-6
- *NEON vectors* on page 9-16.

9.18 Extended notation

The assembler implements an extension to the architectural NEON and VFP assembly syntax, called *extended notation*. This extension allows you to include datatype information or scalar indexes in register names. If you do this, you do not need to include the datatype or scalar index information in every instruction.

Register names can be any of the following:

Untyped The register name specifies the register, but not what datatype it contains, nor any index to a particular scalar within the register.

Untyped with scalar index

The register name specifies the register, but not what datatype it contains, It specifies an index to a particular scalar within the register.

Typed The register name specifies the register, and what datatype it contains, but not any index to a particular scalar within the register.

Typed with scalar index

The register name specifies the register, what datatype it contains, and an index to a particular scalar within the register.

Use the SN, DN, and QN directives to create typed and scalar registers.

9.18.1 See also

Concepts

- *NEON and VFP data types* on page 9-14
- *NEON vectors* on page 9-16
- *NEON scalars* on page 9-22.

Reference

Assembler Reference:

- *QN, DN, and SN* on page 6-16.

9.19 Polynomial arithmetic over $\{0,1\}$

The coefficients 0 and 1 are manipulated using the rules of Boolean arithmetic:

- $0 + 0 = 1 + 1 = 0$
- $0 + 1 = 1 + 0 = 1$
- $0 * 0 = 0 * 1 = 1 * 0 = 0$
- $1 * 1 = 1$.

That is, adding two polynomials over $\{0,1\}$ is the same as a bitwise exclusive OR, and multiplying two polynomials over $\{0,1\}$ is the same as integer multiplication except that partial products are exclusive-ORed instead of being added.

9.19.1 See also

Concepts

- *NEON and VFP data types* on page 9-14.

9.20 NEON and VFP system registers

Three NEON and VFP system registers are accessible in all implementations of NEON and VFP:

- FPSCR, the floating-point status and control register
- FPEXC, the floating-point exception register
- FPSID, the floating-point system ID register.

A particular implementation of NEON or VFP can have additional registers.

9.20.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *FPEXC, the floating-point exception register* on page 9-29
- *FPSID, the floating-point system ID register* on page 9-30
- *Read-Modify-Write procedure* on page 5-33.

Reference

- *Technical Reference Manual* for your VFP Coprocessor.

9.21 FPSCR, the floating-point status and control register

The FPSCR contains all the user-level NEON and VFP status and control bits. NEON only uses bits[31:27]. The bits are used as follows:

bits[31:28] Are the N, Z, C, and V flags. These are the NEON and VFP status flags. They cannot be used to control conditional execution until they have been copied into the status flags in the CPSR.

bit[27] Is the QC, cumulative saturation flag. This is set if saturation occurs in NEON saturating instructions.

bit[25] Is the *Default NaN* (DN) mode control bit:

0 Disabled. NaN operands propagate through to the output of a floating-point operation.

1 Enabled. Any operation involving one or more NaNs returns the Default NaN.

———— **Note** ————

NEON always uses the Default NaN enabled setting regardless of this bit.

bit[24] Is the flush-to-zero mode control bit:

0 Flush-to-zero mode is disabled.

1 Flush-to-zero mode is enabled.

Flush-to-zero mode can provide greater performance, depending on your hardware and software, at the expense of loss of range.

———— **Note** ————

NEON always uses flush-to-zero mode regardless of this bit.

Flush-to-zero mode must not be used when IEEE 754 compatibility is a requirement.

bits[23:22] Control rounding mode as follows:

0b00 *Round to Nearest* (RN) mode.

0b01 *Round towards Plus infinity* (RP) mode.

0b10 *Round towards Minus infinity* (RM) mode.

0b11 *Round towards Zero* (RZ) mode.

———— **Note** ————

NEON always uses the Round to Nearest mode regardless of these bits.

- bits[21:20]** STRIDE is the distance between successive values in a vector. Stride is controlled as follows:
- 0b00** STRIDE = 1
 - 0b11** STRIDE = 2.
- bits[18:16]** LEN is the number of registers used by each vector. It is 1 + the value of bits[18:16]:
- 0b000** LEN = 1
 - ...
 - 0b111** LEN = 8.

bits[15, 12:8] Are the exception trap enable bits:

- IDE** input denormal exception enable
- IXE** inexact exception enable
- UFE** underflow exception enable
- OFE** overflow exception enable
- DZE** division by zero exception enable
- IOE** invalid operation exception enable.

This document does not cover the use of floating-point exception trapping. For information see the technical reference manual for the VFP coprocessor you are using.

bits[7, 4:0] Are the cumulative exception bits:

- IDC** input denormal exception
- IXC** inexact exception
- UFC** underflow exception
- OFC** overflow exception
- DZC** division by zero exception
- IOC** invalid operation exception.

Cumulative exception bits are set when the corresponding exception occurs. They remain set until you clear them by writing directly to the FPSCR.

all other bits Are unused in the basic NEON and VFP specification. They can be used in particular implementations. Do not modify these bits except in accordance with any use in a particular implementation.

To change some bits without affecting other bits, use a read-modify-write procedure.

Note

The use of vector mode is deprecated. Set LEN and STRIDE to 1.

9.21.1 See also

Concepts

- *Flush-to-zero mode* on page 9-31
- *Floating-point exceptions* on page 9-12
- *Conditional execution of NEON and VFP instructions* on page 9-11
- *Condition code meanings* on page 6-9
- *Vectors in the VFP extension register bank* on page 9-36
- *Read-Modify-Write procedure* on page 5-33.

Reference

- The *Technical Reference Manual* for your VFP coprocessor.

9.22 FPEXC, the floating-point exception register

You can only access the FPEXC in privileged software execution. It contains the following bits:

- bit[31]** Is the EX bit. You can read it in all NEON or VFP implementations. In some implementations you might also be able to write to it.
- If the value is 0, the only significant state in the NEON or VFP system is the contents of the general-purpose registers plus FPSCR and FPEXC.
- If the value is 1, you require implementation-specific information to save state.
- bit[30]** Is the EN bit. You can read and write it in all NEON or VFP implementations.
- If the value is 1, NEON (if present) and VFP (if present) are enabled and operate normally.
- If the value is 0, NEON and VFP are disabled. When they are disabled, you can read or write the FPSID or FPEXC registers, but other NEON or VFP instructions are treated as Undefined Instructions.
- bits[29:0]** Might be used by particular implementations of VFP. You can use the VFP instructions without accessing these bits.
- You must not alter these bits except in accordance with their use in a particular implementation.

To change some bits without affecting other bits, use a read-modify-write procedure.

9.22.1 See also

Concepts

- *Flush-to-zero mode* on page 9-31
- *Conditional execution of NEON and VFP instructions* on page 9-11
- *Condition code meanings* on page 6-9
- *Read-Modify-Write procedure* on page 5-33.

Reference

- The *Technical Reference Manual* for your VFP coprocessor.

9.23 FPSID, the floating-point system ID register

The FPSID is a read-only register. You can read it to find out which implementation of the NEON or VFP architecture your program is running on.

bit[31:24]	Implementor code
bit[23]	0 for hardware coprocessor, 1 for software implementation
bit[22:16]	Subarchitecture version number
bit[15:8]	Implementation defined part number
bit[7:4]	Implementation defined variant number
bit[3:0]	Implementation defined revision number.

9.23.1 See also

Concepts

- *NEON and VFP system registers* on page 9-25
- *Read-Modify-Write procedure* on page 5-33.

Reference

- The *Technical Reference Manual* for your VFP coprocessor or floating-point capable processor.

9.24 Flush-to-zero mode

Some implementations of VFP use support code to handle denormalized numbers. The performance of such systems, in calculations involving denormalized numbers, is much less than it is in normal calculations.

Flush-to-zero mode replaces denormalized numbers with 0. This does not comply with IEEE 754 arithmetic, but in some circumstances can improve performance considerably.

NEON and VFPv3 flush-to-zero preserves the sign bit. VFPv2 flush-to-zero flushes to +0.

NEON always uses flush-to-zero mode.

9.24.1 See also

Concepts

- *When to use flush-to-zero mode* on page 9-32
- *The effects of using flush-to-zero mode* on page 9-33
- *Operations not affected by flush-to-zero mode* on page 9-34.

9.25 When to use flush-to-zero mode

You should select flush-to-zero mode if all the following are true:

- IEEE 754 compliance is not a requirement for your system
- the algorithms you are using are such that they sometimes generate denormalized numbers
- your system uses support code to handle denormalized numbers
- the algorithms you are using do not depend for their accuracy on the preservation of denormalized numbers
- the algorithms you are using do not generate frequent exceptions as a result of replacing denormalized numbers with 0.

You can change between flush-to-zero and normal mode at any time, if different parts of your code have different requirements. Numbers already in registers are not affected by changing mode.

9.25.1 See also

Concepts

- *The effects of using flush-to-zero mode* on page 9-33
- *Flush-to-zero mode* on page 9-31.

9.26 The effects of using flush-to-zero mode

With certain exceptions, flush-to-zero mode has the following effects on floating-point operations:

- A denormalized number is treated as 0 when used as an input to a floating-point operation. The source register is not altered.
- If the result of a single-precision floating-point operation, before rounding, is in the range -2^{-126} to $+2^{-126}$, it is replaced by 0.
- If the result of a double-precision floating-point operation, before rounding, is in the range -2^{-1022} to $+2^{-1022}$, it is replaced by 0.

An inexact exception occurs whenever a denormalized number is used as an operand, or a result is flushed to zero. Underflow exceptions do not occur in flush-to-zero mode.

9.26.1 See also

Concepts

- *Operations not affected by flush-to-zero mode* on page 9-34
- *Flush-to-zero mode* on page 9-31.

9.27 Operations not affected by flush-to-zero mode

The following NEON and VFP operations can be carried out on denormalized numbers even in flush-to-zero mode, without flushing the results to zero:

- copy, absolute value, and negate (VMOV, VMVN, V{Q}ABS, and V{Q}NEG)
- duplicate (VDUP)
- swap (VSWP)
- load and store (VLDR and VSTR)
- load multiple and store multiple (VLDM and VSTM)
- transfer between extension registers and ARM general-purpose registers (VMOV).

9.27.1 See also

Concepts

- *The effects of using flush-to-zero mode* on page 9-33
- *Flush-to-zero mode* on page 9-31.

Reference

Assembler Reference:

- *VDUP* on page 4-32
- *VSWP* on page 4-37
- *VLDR and VSTR* on page 4-9
- *VLDM, VSTM, VPOP, and VPUSH* on page 4-11
- *VMOV, VMVN (register)* on page 4-23
- *VABS, VNEG, and VSQRT* on page 4-92
- *V{Q}ABS and V{Q}NEG* on page 4-52
- *VMOV (between two ARM registers and an extension register)* on page 4-13
- *VMOV (between an ARM register and a NEON scalar)* on page 4-15
- *VMOV (between one ARM register and single precision VFP)* on page 4-16.

9.28 VFP vector mode

Usually the VFP core only works on a single register. However, many VFP arithmetic instructions can also operate on vectors of up to eight single-precision or four double-precision numbers, enabling *Single Instruction Multiple Data* (SIMD) parallelism.

In addition, the floating-point load and store instructions have multiple register forms, enabling vectors to be transferred to and from memory easily.

Note

The use of VFP vector mode is deprecated.

9.28.1 See also

Concepts

- *Architecture support for NEON and VFP* on page 9-3
- *Vectors in the VFP extension register bank* on page 9-36.

Reference

- *ARM Architecture Reference Manual*,
<http://infocenter.arm.com/help/topic/com.arm.doc.subset.arch.reference/index.html>.

9.29 Vectors in the VFP extension register bank

For vector operations, the VFP extension register bank can be viewed as a collection of smaller banks. Each of these smaller banks is treated either as a bank of 8 single-precision registers or as a bank of 4 double-precision registers.

In VFPv2, VFPv3-D16, and VFPv3-D16-FP16 the VFP extension register bank can be viewed as a collection of:

- four banks of single-precision registers, s0 to s7, s8 to s15, s16 to s23, and s24 to s31
- four banks of double-precision registers, d0 to d3, d4 to d7, d8 to d11, and d12 to d15
- any combination of single-precision and double-precision banks.

In VFPv3 and VFPv3-FP16, the VFP extension register bank can be viewed as a collection of:

- four banks of single-precision registers, s0 to s7, s8 to s15, s16 to s23, and s24 to s31
- Eight banks of double-precision registers, d0 to d3, d4 to d7, d8 to d11, d12 to d15, d16 to d19, d20 to d23, d24 to d27, and d28 to d31
- any combination of single-precision and double-precision banks.

See Figure 9-2 and Figure 9-3 for further clarification.




Figure 9-2 VFPv2 register banks




Figure 9-3 VFPv3 register banks

A vector, in a VFP instruction, can use up to eight single-precision registers, or four double-precision registers, from the same bank. The number of registers used by a vector is controlled by the LEN bits in the FPSCR.

Note

The value of the LEN bits is not a sufficient condition to perform vector operations using VFP. Whether a VFP operation is scalar, vector or mixed depends on which bank the specified operand and destination registers are in.

A vector can start from any register and wraps around to the beginning of the bank. The first register used by an operand vector is the register that is specified as the operand in the individual VFP instructions. The first register used by the destination vector is the register that is specified as the destination in the individual VFP instructions.

9.29.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *VFP vector wrap-around* on page 9-38
- *VFP vector stride* on page 9-39
- *Restriction on vector length* on page 9-40
- *Control of scalar, vector, and mixed operations* on page 9-41.

9.30 VFP vector wrap-around

If the vector extends beyond the end of a bank, it wraps around to the beginning of the same bank, for example:

- a vector of length 6 starting at s5 is {s5, s6, s7, s0, s1, s2}
- a vector of length 3 starting at s15 is {s15, s8, s9}
- a vector of length 4 starting at s22 is {s22, s23, s16, s17}
- a vector of length 2 starting at d7 is {d7, d4}
- a vector of length 3 starting at d10 is {d10, d11, d8}.

A vector cannot contain registers from more than one bank.

9.30.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *Vectors in the VFP extension register bank* on page 9-36
- *VFP vector stride* on page 9-39
- *Restriction on vector length* on page 9-40.

9.31 VFP vector stride

Vectors can occupy consecutive registers or they can occupy alternate registers. This is controlled by the STRIDE bits in the FPSCR. For example:

- a vector of length 3, stride 2, starting at s1, is {s1, s3, s5}
- a vector of length 4, stride 2, starting at s6, is {s6, s0, s2, s4}
- a vector of length 2, stride 2, starting at d1, is {d1, d3}.
- a vector of length 4, stride 1, starting at d0, is {d0, d1, d2, d3}.

9.31.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *Vectors in the VFP extension register bank* on page 9-36
- *VFP vector wrap-around* on page 9-38
- *Restriction on vector length* on page 9-40.

9.32 Restriction on vector length

A vector cannot use the same register twice. Enabling for vector wrap-around, this means that you cannot have:

- a single-precision vector with length > 4 and stride = 2
- a double-precision vector with length > 4 and stride = 1
- a double-precision vector with length > 2 and stride = 2.

9.32.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *Vectors in the VFP extension register bank* on page 9-36
- *VFP vector wrap-around* on page 9-38.

9.33 Control of scalar, vector, and mixed operations

You can use VFP arithmetic instructions to operate on:

- scalars
- vectors
- scalars and vectors together.

Use the LEN bits in the FPSCR to control the length of vectors. When LEN is 1 all VFP operations are scalar.

When LEN is greater than 1, the VFP operation can be scalar, vector or mixed. The behavior of VFP arithmetic operations depends on which register bank the destination and operand registers are in.

The first bank of registers, s0 to s7 or d0 to d3 and the fifth bank of registers d16 to d19 are scalar banks. All other banks are vector banks. A vector operation or mixed operation is one where the destination register is in one of the vector banks.

Given instructions of the following general forms:

$$\begin{array}{l} Op \quad Fd, Fn, Fm \\ Op \quad Fd, Fm \end{array}$$

where:

<i>Op</i>	is the VFP instruction
<i>Fd</i>	is the destination register
<i>Fn</i>	is an operand register
<i>Fm</i>	is the only or second operand register

the behavior of the operation is as follows:

- If *Fd* is in the first or fifth bank of registers then the operation is scalar.
- If *Fm* is in the first or fifth bank of registers, but *Fd* is not, then the operation is mixed.
- If neither *Fd* nor *Fm* are in the first or fifth bank of registers, the operation is vector.

In scalar operations, *Op* acts on the value in *Fm*, and the value in *Fn* if present. The result is placed in *Fd*.

In vector operations, *Op* acts on the values in the vector starting at *Fm*, together with the values in the vector starting at *Fn* if present. The results are placed in the vector starting at *Fd*.

In mixed operations, with a single operand, *Op* acts on the single value in *Fm* and LEN copies of the result are placed in the vector starting at *Fd*.

In mixed operations, with two operands, Op acts on the single value in Fm , together with the values in the vector starting at Fn . The results are placed in the vector starting at Fd .

9.33.1 See also

Concepts

- *FPSCR, the floating-point status and control register* on page 9-26
- *Vectors in the VFP extension register bank* on page 9-36
- *VFP vector wrap-around* on page 9-38
- *VFP vector stride* on page 9-39
- *Restriction on vector length* on page 9-40.

9.34 VFP directives and vector notation

This applies only to `armasm`. The inline assemblers in the C and C++ compilers do not accept these directives or vector notation.

The use of VFP vector mode is deprecated, and vector notation is not supported in UAL. To use vector notation, you must use the pre-UAL VFP mnemonics. See *Pre-UAL VFP mnemonics* on page 9-44 for details. You can mix pre-UAL VFP mnemonics and UAL VFP mnemonics.

You can make assertions about VFP vector lengths and strides in your code, and have them checked by the assembler. See:

- *VFPASSERT SCALAR* on page 9-47
- *VFPASSERT VECTOR* on page 9-48.

If you use `VFPASSERT` directives, you must specify vector details in all VFP data processing instructions written using pre-UAL mnemonics. The vector notation is described in *Vector notation* on page 9-46. If you do not use `VFPASSERT` directives you must not use this notation.

9.35 Pre-UAL VFP mnemonics

Where UAL mnemonics use .F32 to specify single-precision data, pre-UAL mnemonics use S appended to the instruction mnemonic. For example, VABS.F32 was FABSS.

Where UAL mnemonics use .F64 to specify double-precision data, pre-UAL mnemonics use D appended to the instruction mnemonic. For example, VCMPE.F64 was FCMPE.D.

Table 9-4 shows the pre-UAL mnemonics of those instructions that are affected by VFP vector mode. All other VFP instructions are always scalar regardless of the settings of LEN and STRIDE.

Table 9-4 Pre-UAL VFP mnemonics

UAL mnemonic	Equivalent pre-UAL mnemonic
VABS	FABS
VADD	FADD
VDIV	FDIV
VMLA	FMAC
VMLS	FNMAC
VMOV (immediate)	FCONST ^a
VMOV (register)	FCPY
VMUL	FMUL
VNEG	FNEG
VNMLA	FNMSC
VNMLS	FMSC
VNMUL	FNMUL
VSQRT	FSQRT
VSUB	FSUB

a. The immediate in VMOV (immediate) is the floating-point number you want to load. The immediate in FCONST is the number encoded in the instruction to produce the floating-point number you want to load. See *Immediate values in FCONST* on page 9-45 for details.

9.35.1 Immediate values in FCONST

Table 9-5 shows the floating-point values you can load using FCONST. Trailing zeroes are omitted for clarity. The immediate value you must put in the FCONST instruction is the decimal representation of the binary number abcdefgh, where:

- a is 0 for positive numbers, or 1 for negative numbers
- bcd is shown in the column headings
- efgh is shown in the row headings.

Alternatively, you can use 0x followed by the hexadecimal representation.

Table 9-5 Floating-point values for use with FCONST

	bcd	000	001	010	011	100	101	110	111
efgh									
0000		2.0	4.0	8.0	16.0	0.125	0.25	0.5	1.0
0001		2.125	4.25	8.5	17.0	0.1328125	0.265625	0.53125	1.0625
0010		2.25	4.5	9.0	18.0	0.140625	0.28125	0.5625	1.125
0011		2.375	4.75	9.5	19.0	0.1484375	0.296875	0.59375	1.1875
0100		2.5	5.0	10.0	20.0	0.15625	0.3125	0.625	1.25
0101		2.625	5.25	10.5	21.0	0.1640625	0.328125	0.65625	1.3125
0110		2.75	5.5	11.0	22.0	0.171875	0.34375	0.6875	1.375
0111		2.875	5.75	11.5	23.0	0.1796875	0.359375	0.71875	1.4375
1000		3.0	6.0	12.0	24.0	0.1875	0.375	0.75	1.5
1001		3.125	6.25	12.5	25.0	0.1953125	0.390625	0.78125	1.5625
1010		3.25	6.5	13.0	26.0	0.203125	0.40625	0.8125	1.625
1011		3.375	6.75	13.5	27.0	0.2109375	0.421875	0.84375	1.6875
1100		3.5	7.0	14.0	28.0	0.21875	0.4375	0.875	1.75
1101		3.625	7.25	14.5	29.0	0.2265625	0.453125	0.90625	1.8125
1110		3.75	7.5	15.0	30.0	0.234375	0.46875	0.9375	1.875
1111		3.875	7.75	15.5	31.0	0.2421875	0.484375	0.96875	1.9375

9.36 Vector notation

In pre-UAL VFP data processing instructions, specify vectors of VFP registers using angle brackets:

- sn is a single-precision scalar register n .
- $sn\langle\rangle$ is a single-precision vector whose length and stride are given by the current vector length and stride, as defined by `VFPASSERT VECTOR`. The vector starts at register n .
- $sn\langle L\rangle$ is a single-precision vector of length L , stride 1. The vector starts at register n .
- $sn\langle L:S\rangle$ is a single-precision vector of length L , stride S . The vector starts at register n .
- dn is a double-precision scalar register n .
- $dn\langle\rangle$ is a double-precision vector whose length and stride are given by the current vector length and stride, as defined by `VFPASSERT VECTOR`. The vector starts at register n .
- $dn\langle L\rangle$ is a double-precision vector of length L , stride 1. The vector starts at register n .
- $dn\langle L:S\rangle$ is a double-precision vector of length L , stride S . The vector starts at register n .

You can use this vector notation with names defined using the `DN` and `SN` directives.

You must not use this vector notation in the `DN` and `SN` directives themselves.

9.36.1 See also

Concepts

- `VFPASSERT VECTOR` on page 9-48.

Reference

Assembler Reference:

- `QN`, `DN`, and `SN` on page 6-16.

9.37 VFPASSERT SCALAR

The VFPASSERT SCALAR directive informs the assembler that following VFP instructions are in scalar mode. This forces the instruction syntax to be scalar.

9.37.1 Syntax

VFPASSERT SCALAR

9.37.2 Usage

Use the VFPASSERT SCALAR directive to mark the end of any block of code where the VFP mode is VECTOR.

Place the VFPASSERT SCALAR directive immediately after the instruction where the change occurs. This is usually an FMXR instruction, but might be a BL instruction.

If a function expects the VFP to be in vector mode on exit, place a VFPASSERT SCALAR directive immediately after the last instruction. Such a function would not be AAPCS compliant.

————— Note —————

This directive does not generate any code. It is only an assertion by the programmer. The assembler produces error messages if any such assertions are inconsistent with each other, or with any vector notation in VFP data processing instructions.

The assembler faults vector notation in VFP data processing instructions following a VFPASSERT SCALAR directive, even if the vector length is 1.

9.37.3 Example

```
VFPASSERT SCALAR ; scalar mode
fadd d4, d4, d0 ; okay
fadd s4<3>, s0, s8<3> ; ERROR, vector in scalar mode
fabss s24<1>, s28<1> ; ERROR, vector in scalar mode
; (even though length==1)
```

9.37.4 See also

Reference

- *Vector notation* on page 9-46
- *VFPASSERT VECTOR* on page 9-48
- *Procedure Call Standard for the ARM Architecture*,
<http://infocenter.arm.com/help/topic/com.arm.doc.ih0042-/index.html>.

9.38 VFPASSERT VECTOR

The VFPASSERT VECTOR directive informs the assembler that following VFP instructions are in vector mode. It can also specify the length and stride of the vectors.

9.38.1 Syntax

```
VFPASSERT VECTOR{<{n:s}>>}
```

where:

n is the vector length, 1-8.

s is the vector stride, 1-2.

9.38.2 Usage

Use the VFPASSERT VECTOR directive to mark the start of a block of instructions where the VFP mode is VECTOR, and to mark changes in the length or stride of vectors.

Place the VFPASSERT VECTOR directive immediately after the instruction where the change occurs. This is usually an FMXR instruction, but might be a BL instruction.

If a function expects the VFP to be in vector mode on entry, place a VFPASSERT VECTOR directive immediately before the first instruction. Such a function would not be AAPCS compliant.

———— Note ————

This directive does not generate any code. It is only an assertion by the programmer. The assembler produces error messages if any such assertions are inconsistent with each other, or with any vector notation in VFP data processing instructions.

9.38.3 Example

```

VMRS    r10,FPSCR           ; UAL mnemonic - could be FMRX instead.
BIC     r10,r10,#0x00370000
ORR     r10,r10,#0x00020000   ; set length = 3, stride = 1
VMSR    FPSCR,r10
VFPASSERT VECTOR           ; assert vector mode, unspecified length & stride
fadddd  d4, d4, d0           ; ERROR, scalar in vector mode
fadds   s16<3>, s0, s8<3>     ; okay
fabss   s24<1>, s28<1>       ; wrong length, but not faulted (unspecified)
VMRS    r10,FPSCR
BIC     r10,r10,#0x00370000
ORR     r10,r10,#0x00030000   ; set length = 4, stride = 1
VMSR    FPSCR,r10
VFPASSERT VECTOR<4>        ; assert vector mode, length 4, stride 1
fadds   s24<4>, s0, s8<4>     ; okay
fabss   s24<2>, s24<2>       ; ERROR, wrong length
VMRS    r10,FPSCR
BIC     r10,r10,#0x00370000
ORR     r10,r10,#0x00130000   ; set length = 4, stride = 2
VMSR    FPSCR,r10
VFPASSERT VECTOR<4:2>      ; assert vector mode, length 4, stride 2
fadds   s8<4>, s0, s16<4>     ; ERROR, wrong stride because omitting the stride
                                   ; causes a default stride of 1.
fabss   s16<4:2>, s28<4:2>    ; okay
fadds   s8<>, s2, s16<>       ; okay (s8 and s16 both have
                                   ; length 4 and stride 2.
                                   ; s2 is scalar.)

```

9.38.4 See also

Concepts

- *Vector notation* on page 9-46.

Reference

- *Vector notation* on page 9-46
- *VFPASSERT SCALAR* on page 9-47
- *Procedure Call Standard for the ARM Architecture*,
<http://infocenter.arm.com/help/topic/com.arm.doc.ih0042-/index.html>.

